

Yandex



CatBoost: Gradient Boosting for data with both numerical and text features

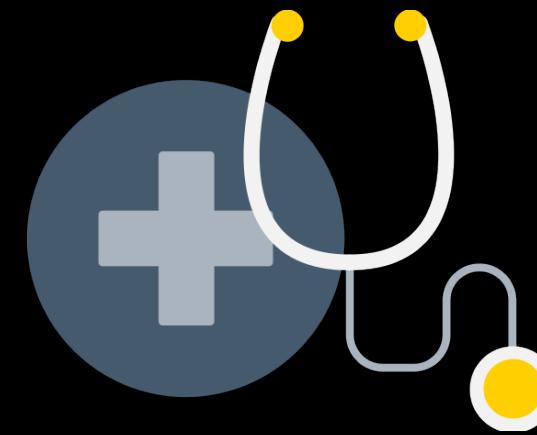
Kirillov Stanislav,
Head of CatBoost team

Plan

- › Gradient boosting and CatBoost overview
- › CatBoost in Yandex and in the wild
- › Supported feature types and text support
- › Big data support and other new features

Applied ML (supervised learning)

Applications



Medicine

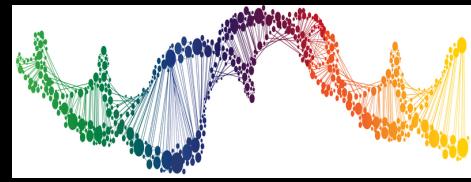


Finance



Sales prediction

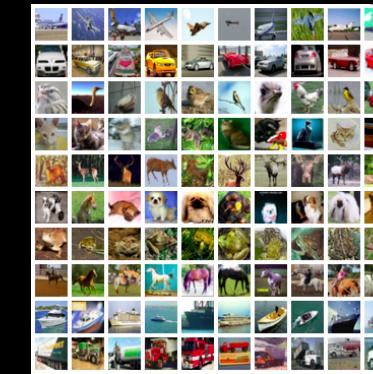
Data at hand



DNA



Text



Images



Music

Tool

› Linear models

› Neural nets

› Decision trees

› GBDT

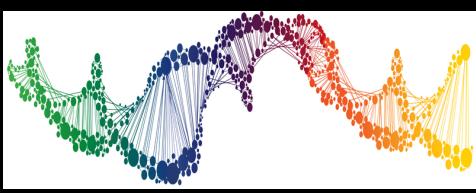
› etc

Data at hand

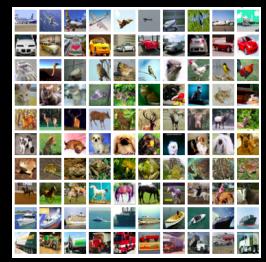
Unstructured data



Music



DNA



Images



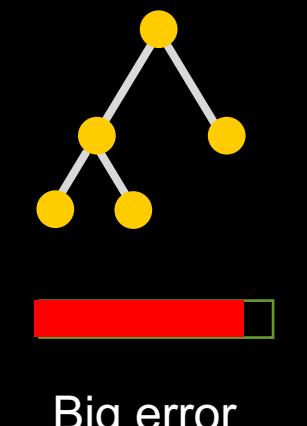
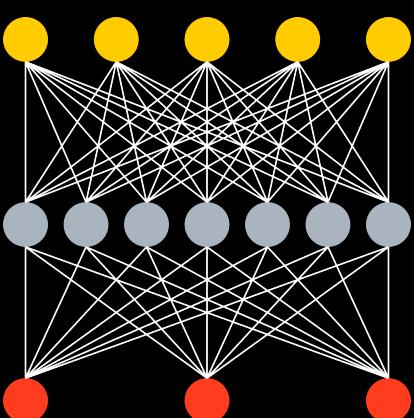
Text

Tabular (or structured) data

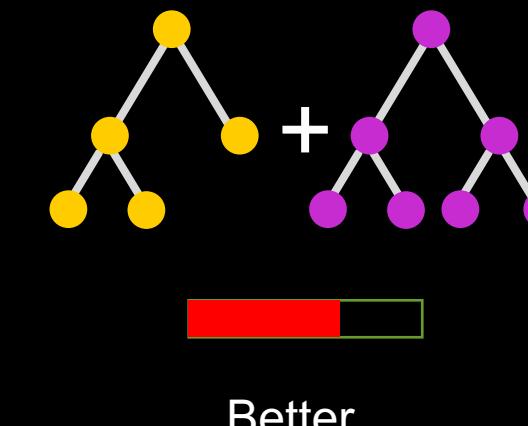
Well engineered features

Music track length	Year	Rating	Label
2	1990	3	1
3	1950	5	0
15	1970	4	1

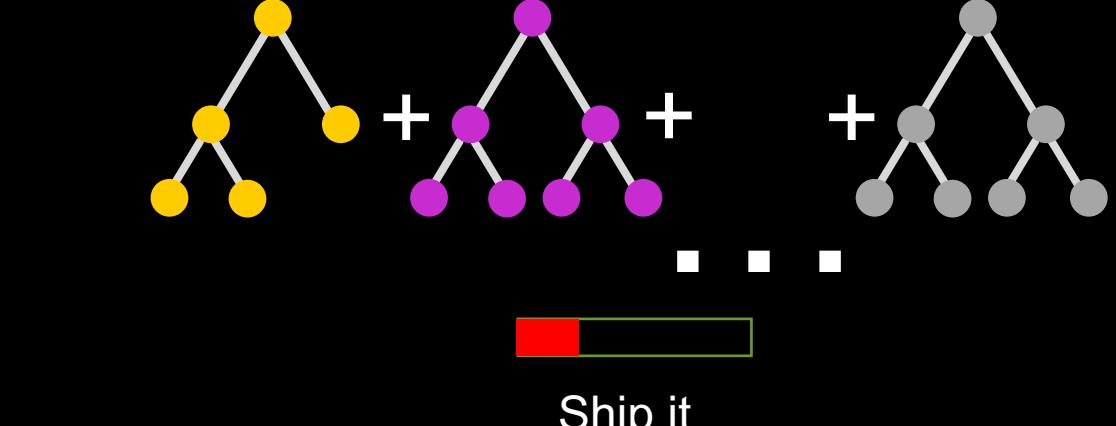
↓
End2End with Deep NN



Big error



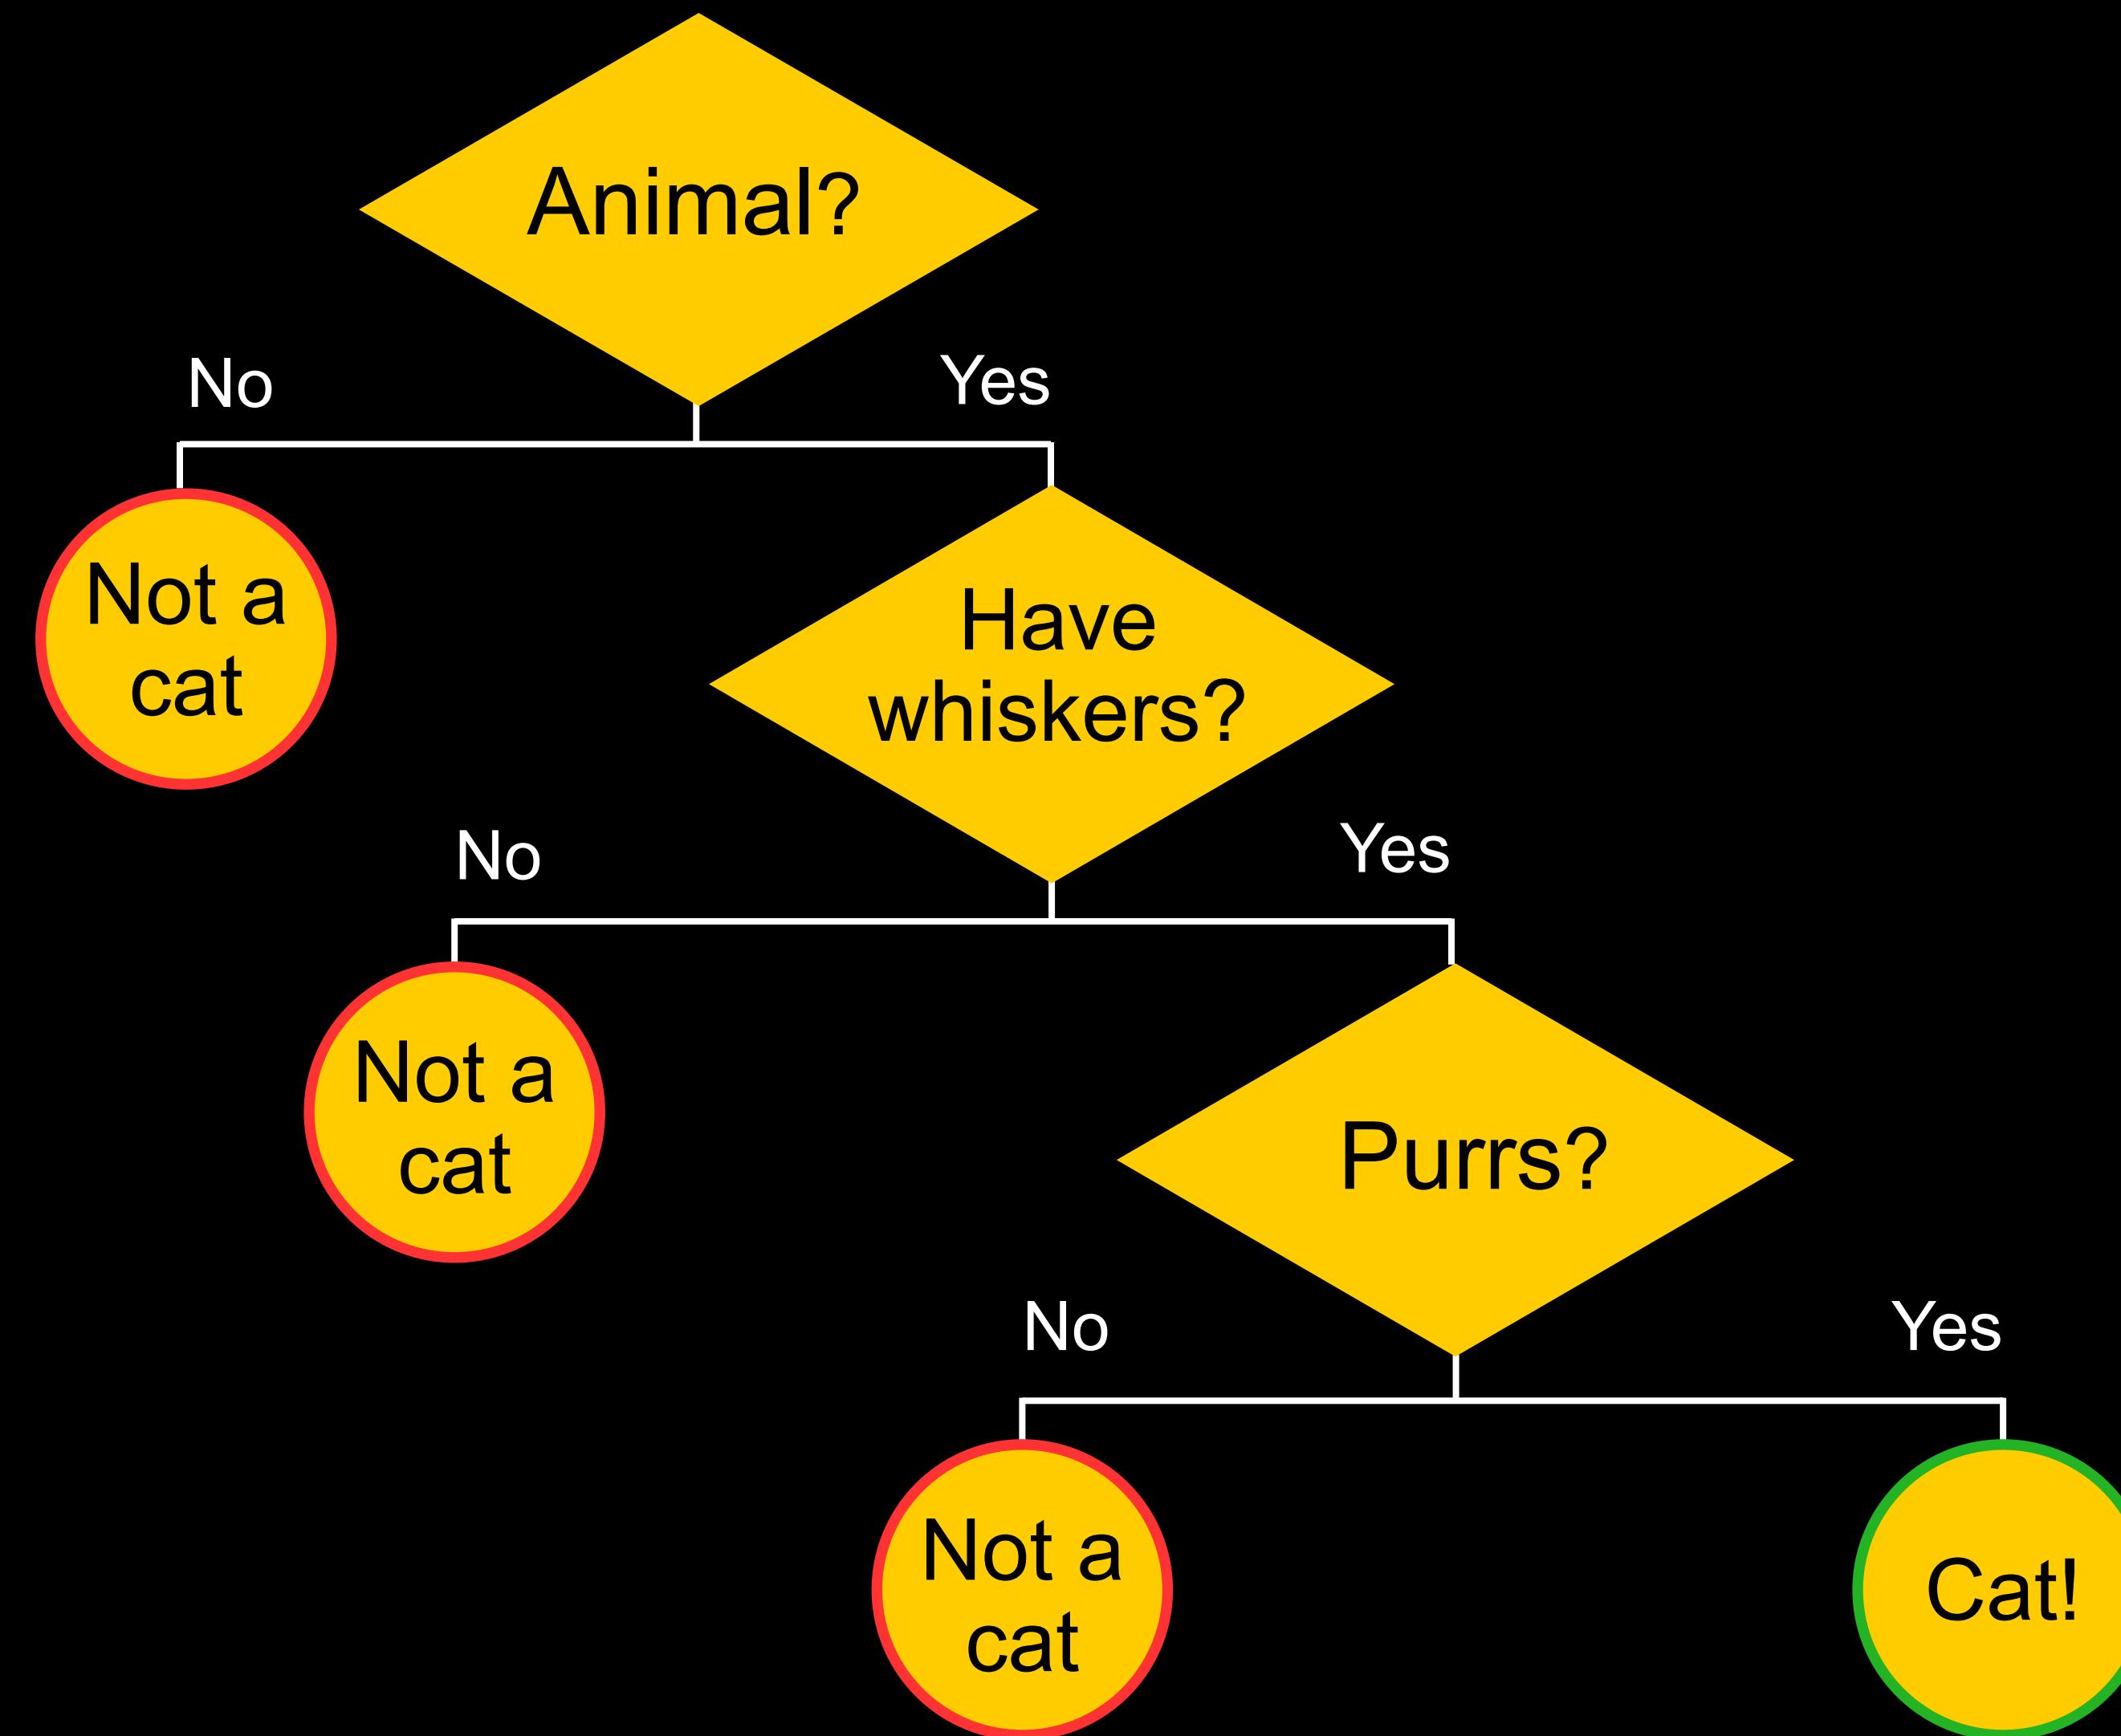
Better



Ship it

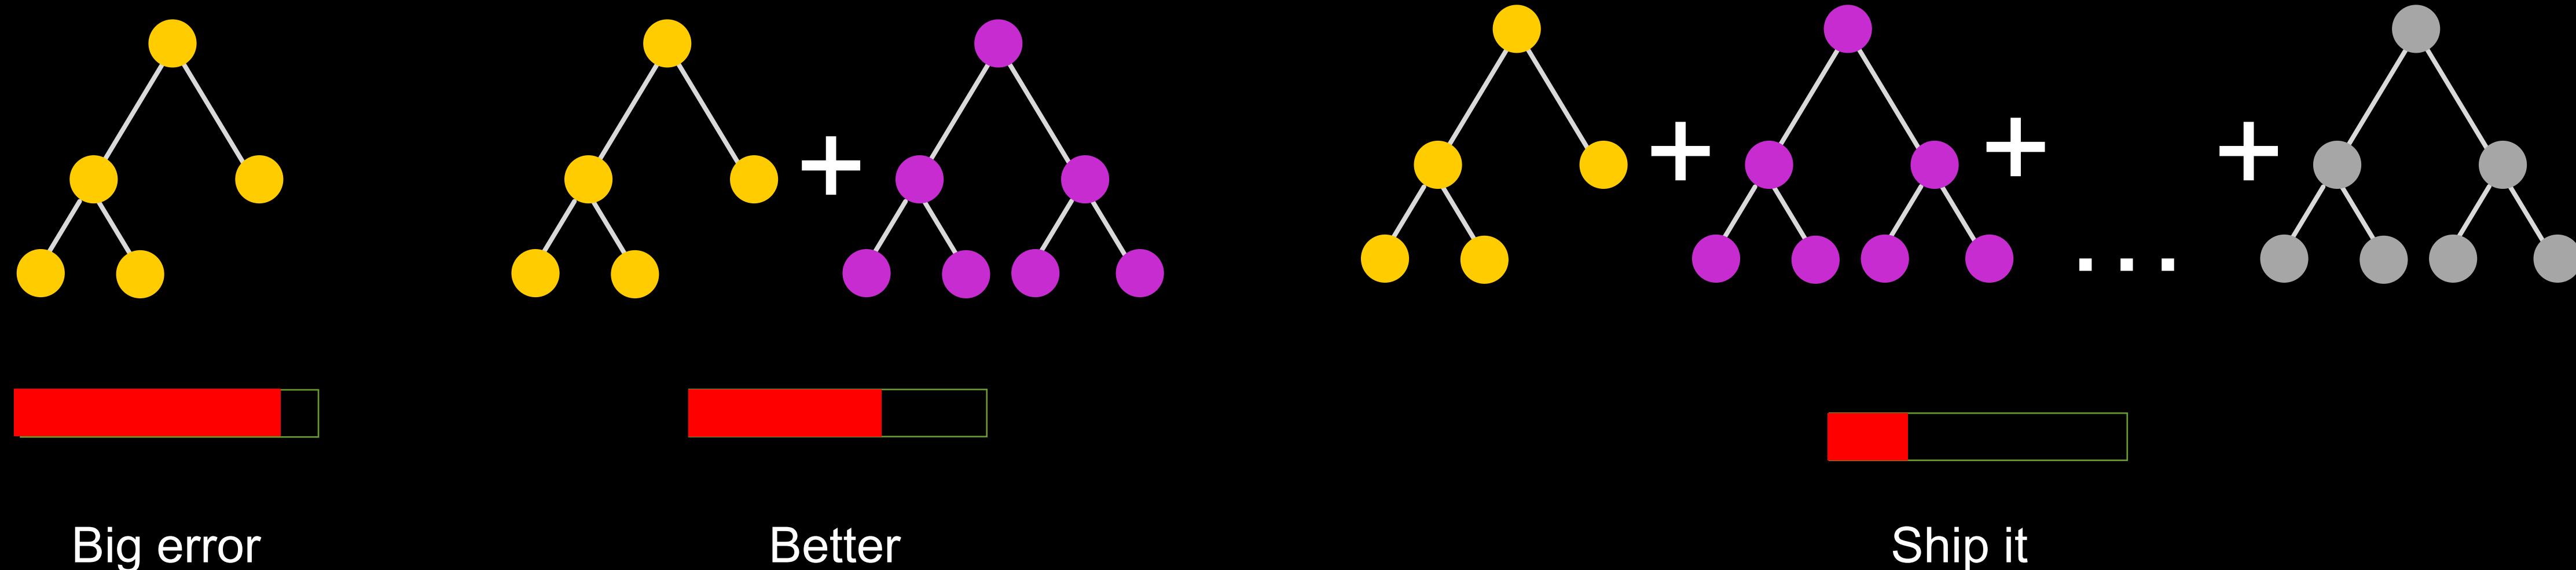
↓
GBDT

Tabular data? Decision trees!

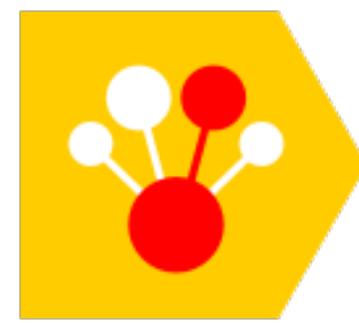


Gradient boosted decision trees

- | State-of-the-art quality on tabular data
- | Easy to use, no sophisticated parameter tuning
- | Works well with small data and scales for big data problems



Main Boosting libraries



Yandex
CatBoost



Microsoft
LightGBM

CatBoost

- 50K pip installs per week
- 5K stars on github
- 64 releases

The screenshot shows the GitHub repository page for `catboost / catboost`. The top right corner shows 5k stars, with a red heart drawn over the star icon. Below the star count, the repository has 189 issues, 7 pull requests, 0 actions, 0 security, 189 insights, and settings. The repository description states: "A fast, scalable, high performance Gradient Boosting on Decision Trees library, used for ranking, classification, regression and other machine learning tasks for Python, R, Java, C++. Supports computation on CPU and GPU. <https://catboost.ai>". The repository has 10,476 commits, 12 branches, 0 packages, 64 releases (highlighted with a red heart), 171 contributors, and is licensed under Apache-2.0. The commit list shows the following recent changes:

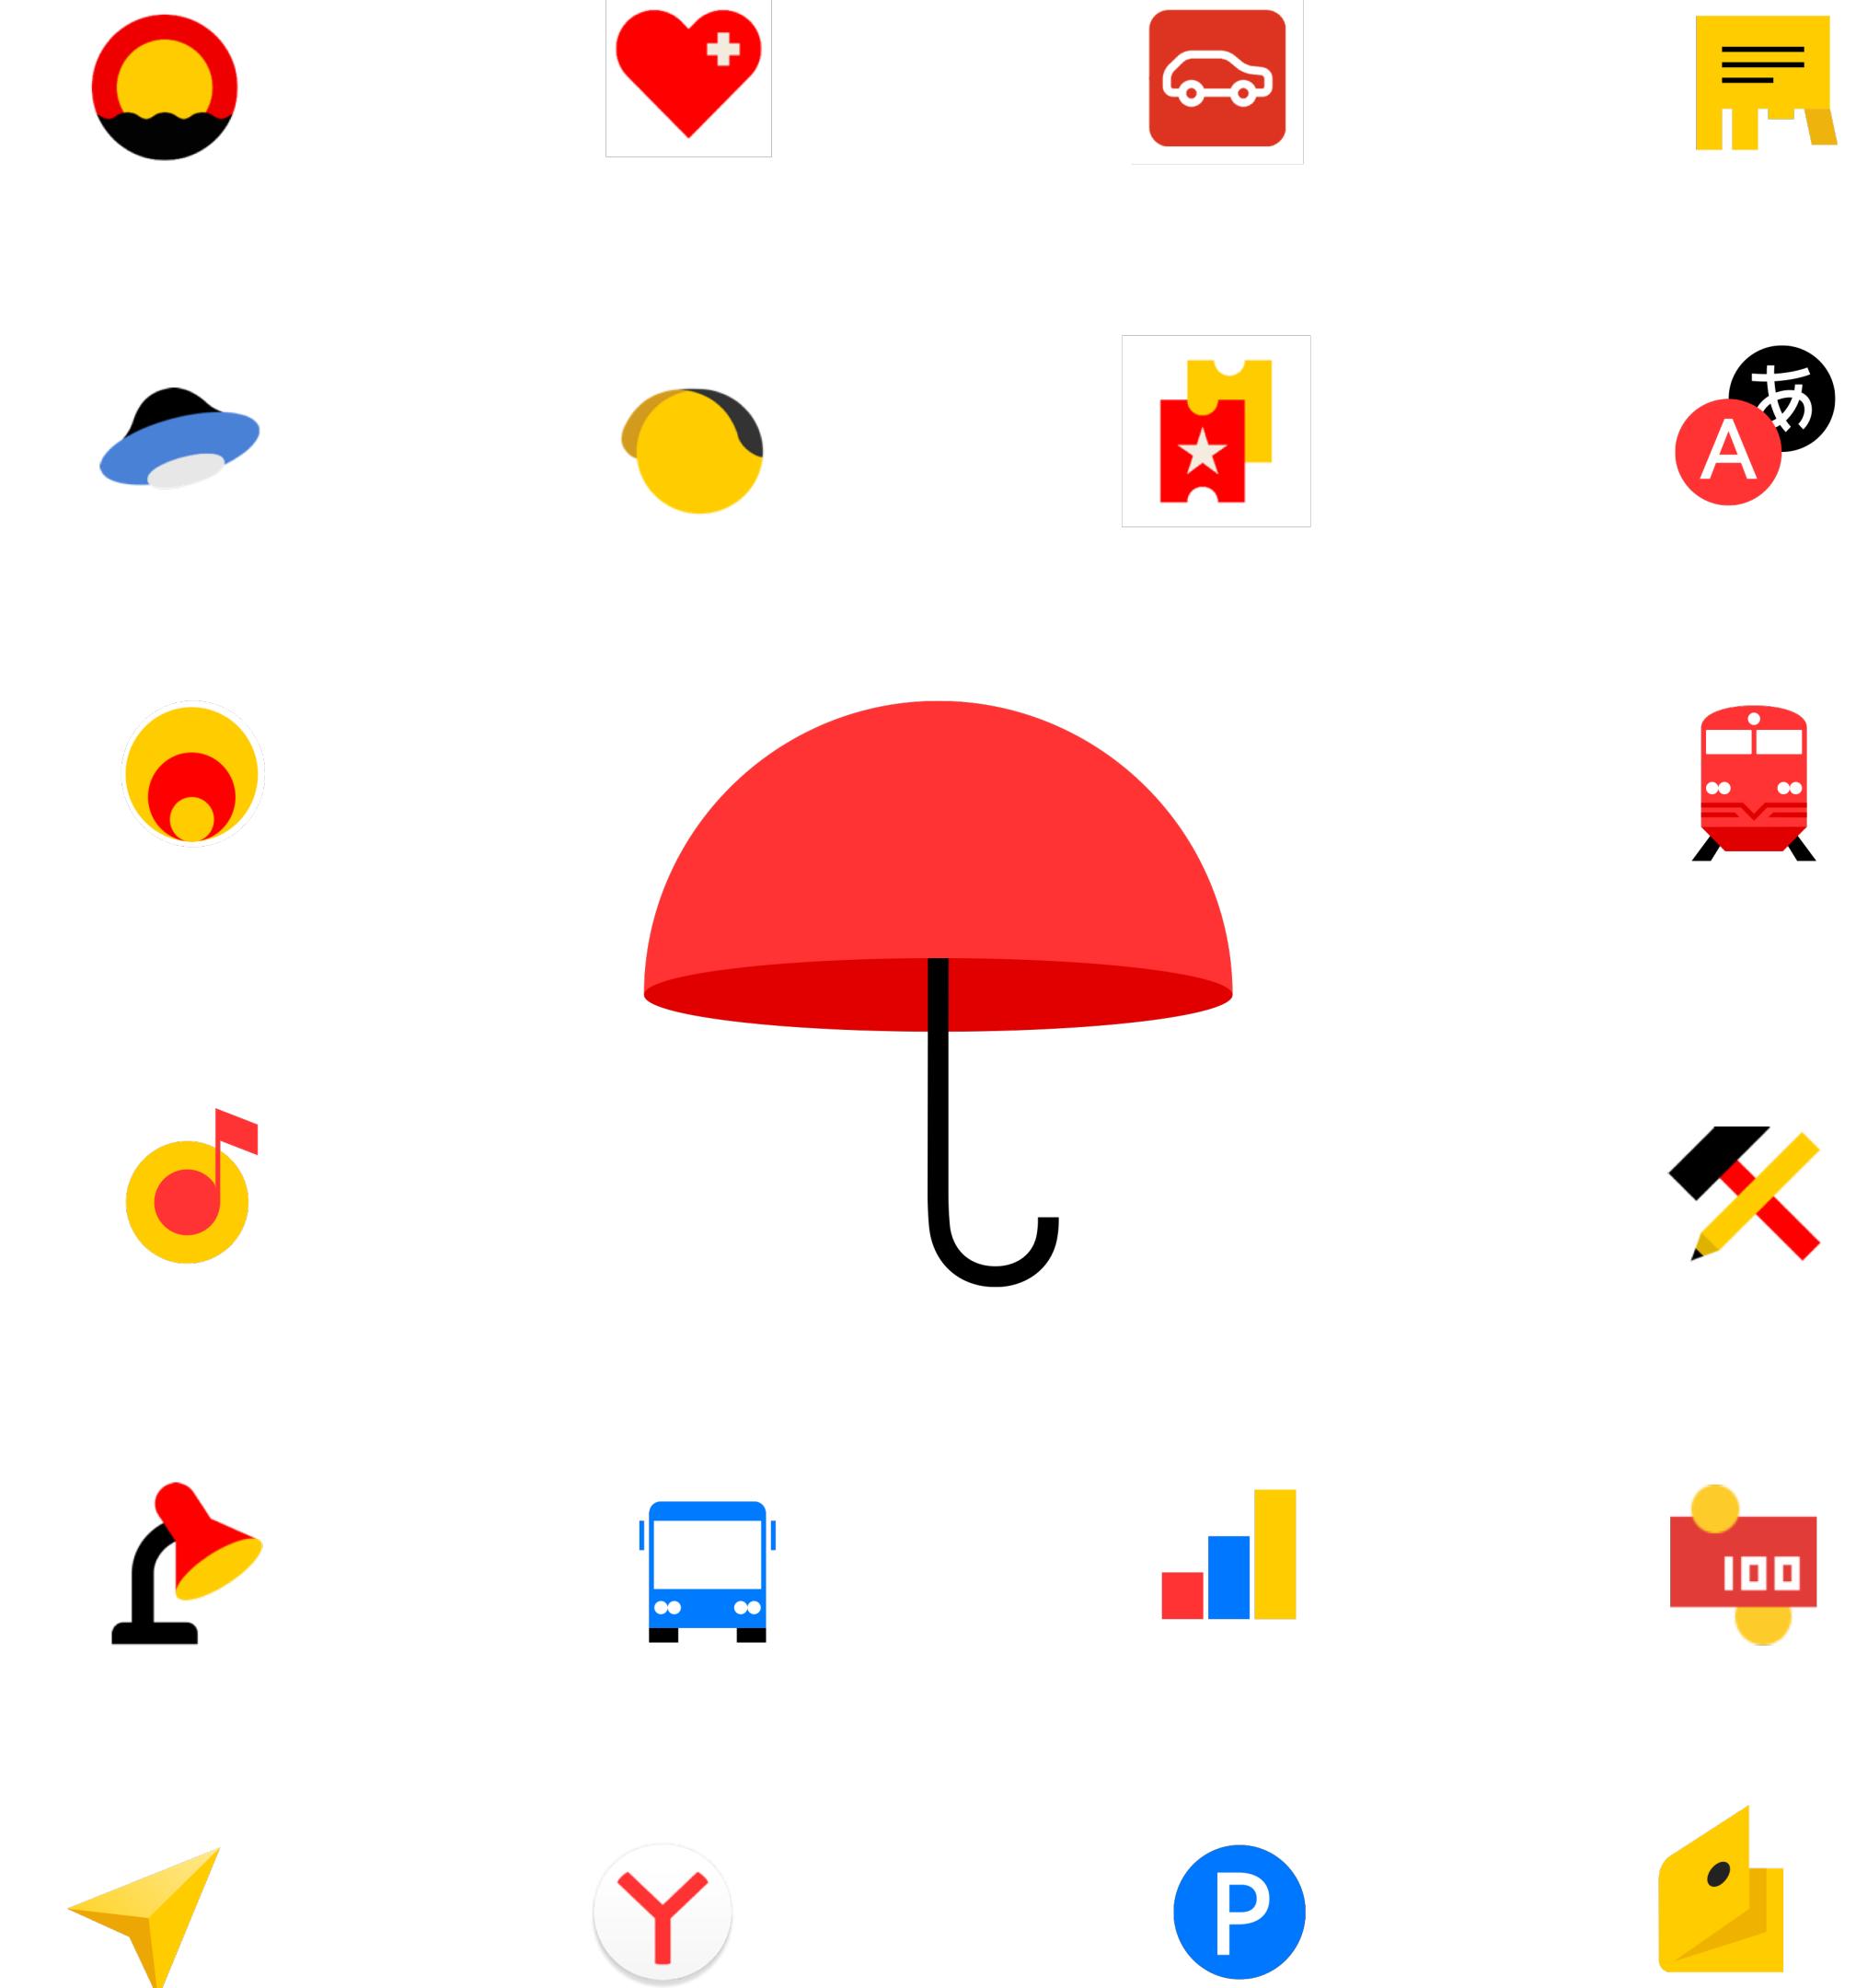
Commit	Author	Message	Time
fec53d1	borman	CONTRIB-1618: Switch metrika/ code to new glib version	Latest commit 3 hours ago
		CONTRIB-1618: Switch metrika/ code to new glib version	3 hours ago
		Force Intel MKL to use SSE4.2 ISA in pytest	3 hours ago
		intermediate changes	10 months ago
		remove custom cuda compiler hack-string from build_all_win	7 days ago
		Specify NO_RUNTIME / NO_UTIL for contrib/libs/clapack	5 hours ago
		intermediate changes	11 hours ago
		intermediate changes	9 months ago
		Solution files and make files updated	5 hours ago
		Solution files and make files updated	5 hours ago

CatBoost advantages

- › Good quality with default parameters
- › Sophisticated categorical and text features support
- › Model analysis tools
- › Set of tools to make GBDT usage easier

CatBoost in Yandex

- Yandex.Zen
- Yandex.Music
- Yandex.Self-Driving Cars
- Yandex.Search
- Yandex.Ads
- Yandex.Weather
- Yandex Alice
- Practically everywhere!



Yandex.Search

Task?

- › Search document order prediction

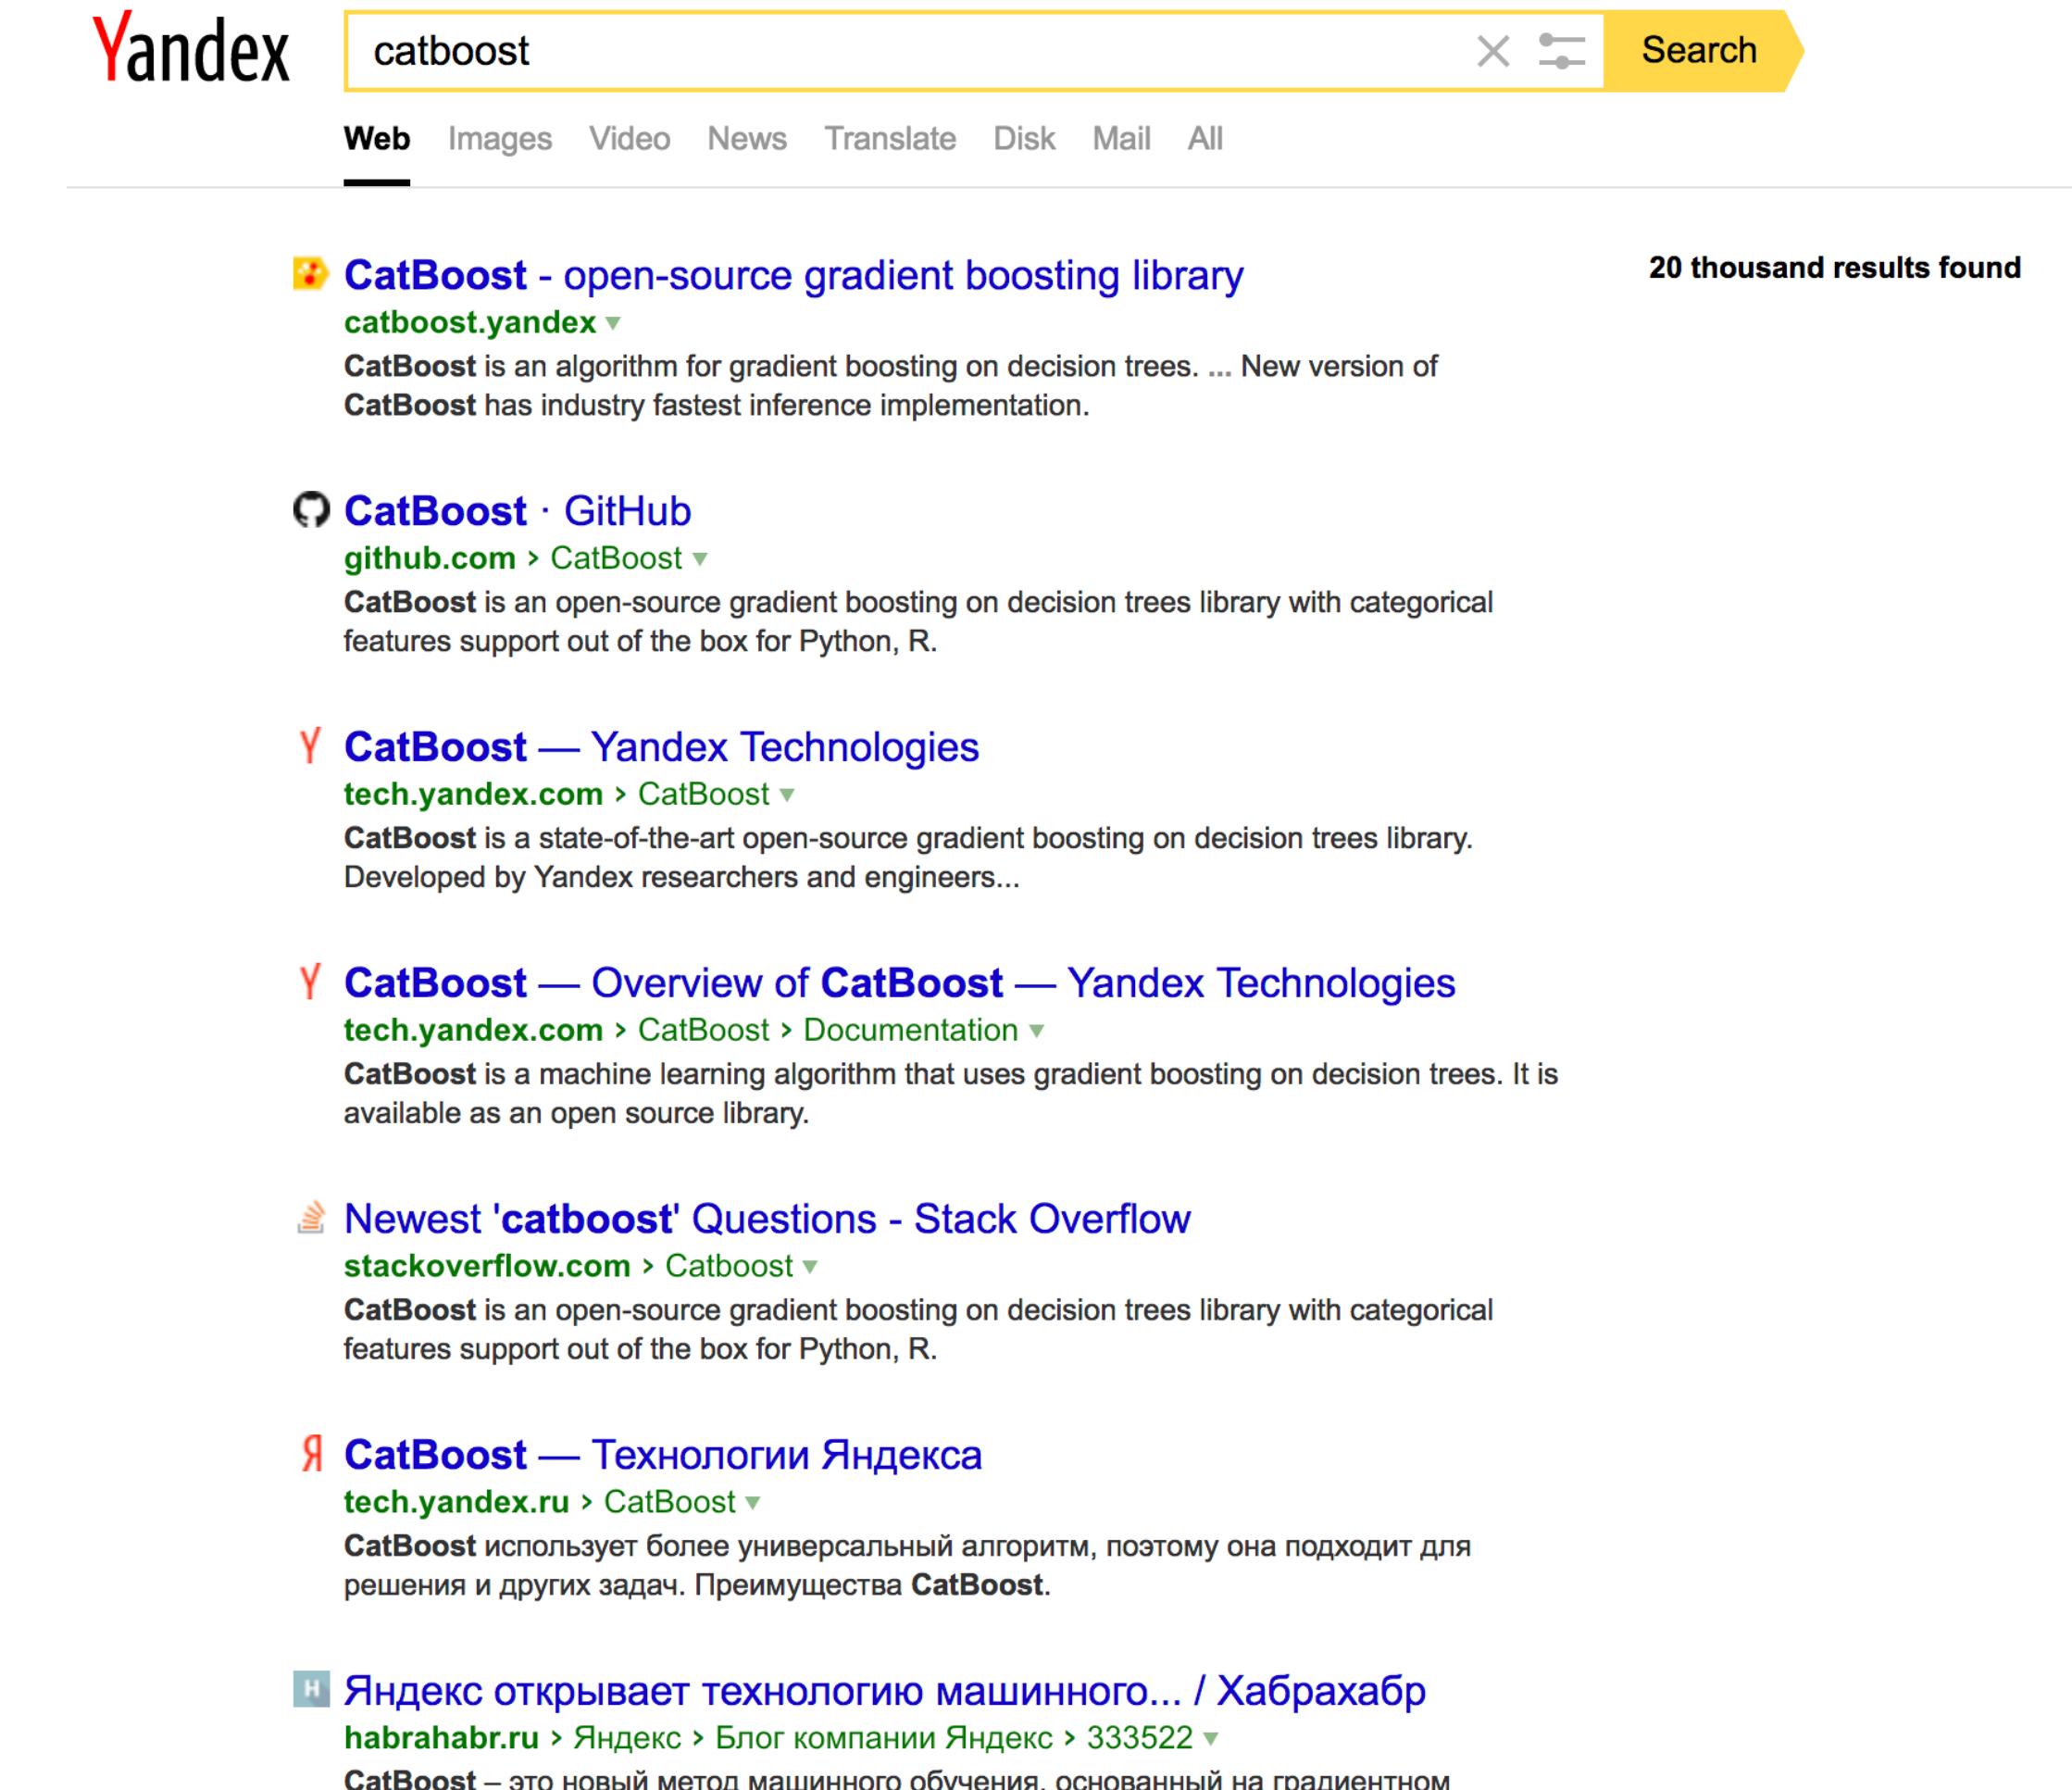
Task type: ranking

Dataset features:

- › Classic features (PageRank, BM25 and others)
- › Neural Networks output

CatBoost features used:

- › YetiRankPairwise target
- › Distributed GPU training
- › Model blending
- › Feature importance analysis
- › Ranking analysis



The screenshot shows the Yandex search results for the query "catboost". The search bar at the top contains "catboost". Below the search bar, a navigation bar includes "Web", "Images", "Video", "News", "Translate", "Disk", "Mail", and "All". The results section displays 20 thousand results found. The first result is a link to "CatBoost - open-source gradient boosting library" from "catboost.yandex". The second result is a link to "CatBoost · GitHub" from "github.com". The third result is a link to "CatBoost — Yandex Technologies" from "tech.yandex.com". The fourth result is a link to "CatBoost — Overview of CatBoost — Yandex Technologies" from "tech.yandex.com". The fifth result is a link to "Newest 'catboost' Questions - Stack Overflow" from "stackoverflow.com". The sixth result is a link to "CatBoost — Технологии Яндекса" from "tech.yandex.ru". The seventh result is a link to "Яндекс открывает технологию машинного... / Хабрахабр" from "habrahabr.ru". Each result includes a snippet of text describing the content of the page.

catboost

Web Images Video News Translate Disk Mail All

20 thousand results found

CatBoost - open-source gradient boosting library
catboost.yandex

CatBoost is an algorithm for gradient boosting on decision trees. ... New version of CatBoost has industry fastest inference implementation.

CatBoost · GitHub
github.com > CatBoost

CatBoost is an open-source gradient boosting on decision trees library with categorical features support out of the box for Python, R.

CatBoost — Yandex Technologies
tech.yandex.com > CatBoost

CatBoost is a state-of-the-art open-source gradient boosting on decision trees library. Developed by Yandex researchers and engineers...

CatBoost — Overview of CatBoost — Yandex Technologies
tech.yandex.com > CatBoost > Documentation

CatBoost is a machine learning algorithm that uses gradient boosting on decision trees. It is available as an open source library.

Newest 'catboost' Questions - Stack Overflow
stackoverflow.com > Catboost

CatBoost is an open-source gradient boosting on decision trees library with categorical features support out of the box for Python, R.

CatBoost — Технологии Яндекса
tech.yandex.ru > CatBoost

CatBoost использует более универсальный алгоритм, поэтому она подходит для решения и других задач. Преимущества CatBoost.

Яндекс открывает технологию машинного... / Хабрахабр
habrahabr.ru > Яндекс > Блог компании Яндекс > 333522

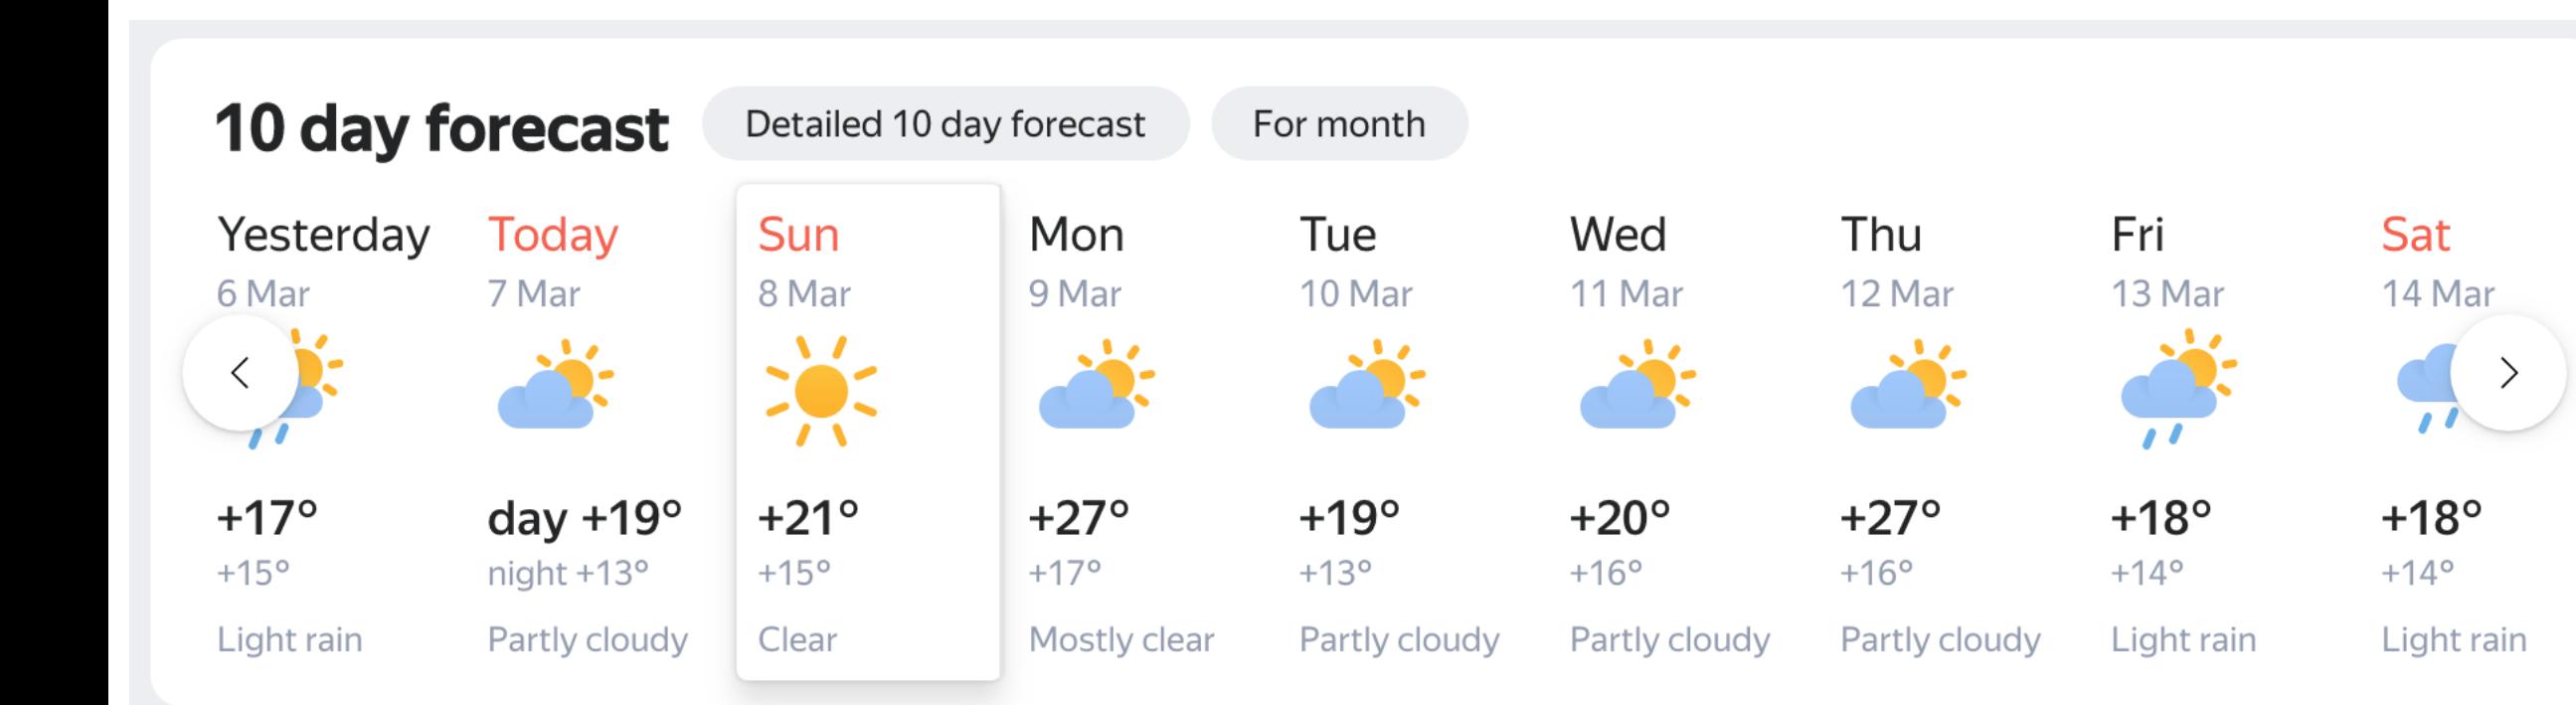
CatBoost – это новый метод машинного обучения, основанный на градиентном

Yandex.Weather

- | **Task?**
- | > Cloudiness type and temperature prediction
- | **Task type: multiclassification and regression**

- | **Dataset features**
- | > Physical weather model output
- | > Neural network output
- | > Online-data from weather stations
- | > Weather historical data

- | **CatBoost features used:**
- | > Multiclassification target and RMSE (for temperature)
- | > GPU training
- | > Feature importance analysis
- | > Training process visualization



Yandex.Alice

Task?

- › Intent classification
- › Select answer in chit chat mode

Task type: classification, ranking

Dataset features:

- › Features based on dialog context
- › Features based on suggested answer

CatBoost features used:

- › GPU training
- › Feature importance analysis
- › Training process visualization



Смотрела фильм "она"?

Have you seen the movie "Her"?

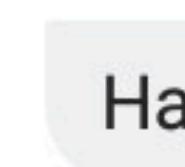


Он очень страшный?

Is it very scary?

Он очень добрый и хороший

It is very kind and nice

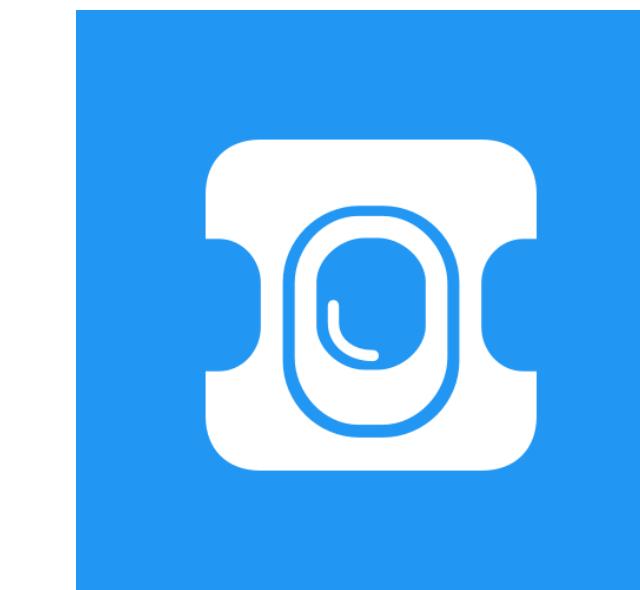


Надо будет посмотреть

I should watch it then

CatBoost in the Wild

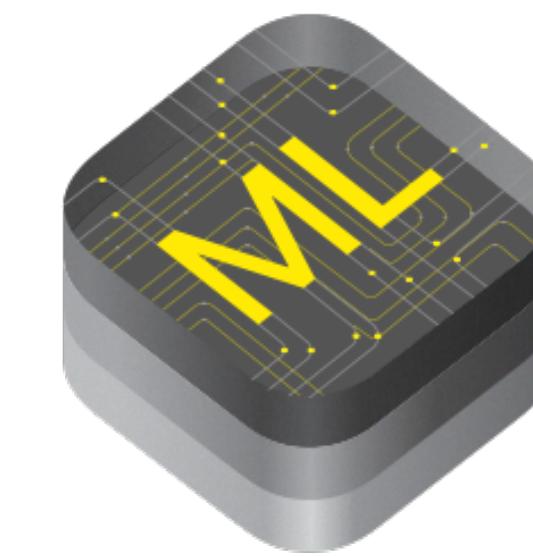
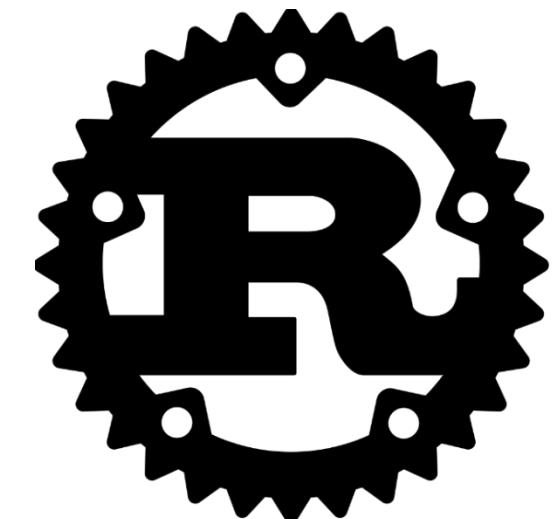
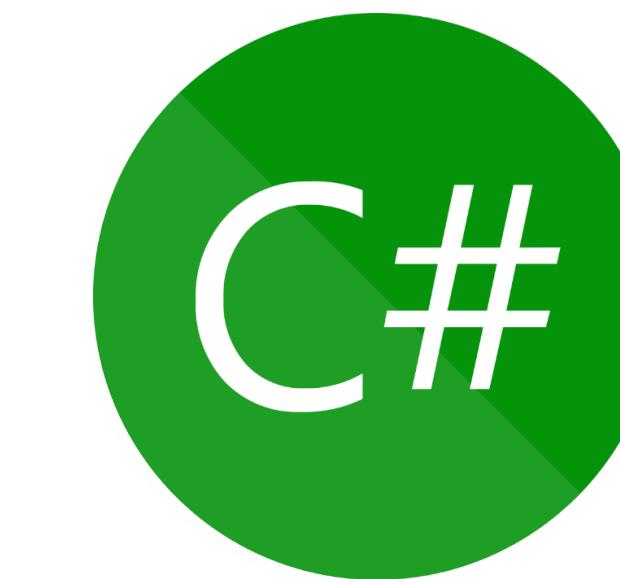
- › Recommendations at Netflix
- › Hotel ranking in Aviasales
- › Protection against bots in CloudFlare
- › Particle classification in CERN
- › Medical research at University of NSW Sydney
- › Destination prediction in Careem taxi service
- › ML competitions on Kaggle



kaggle

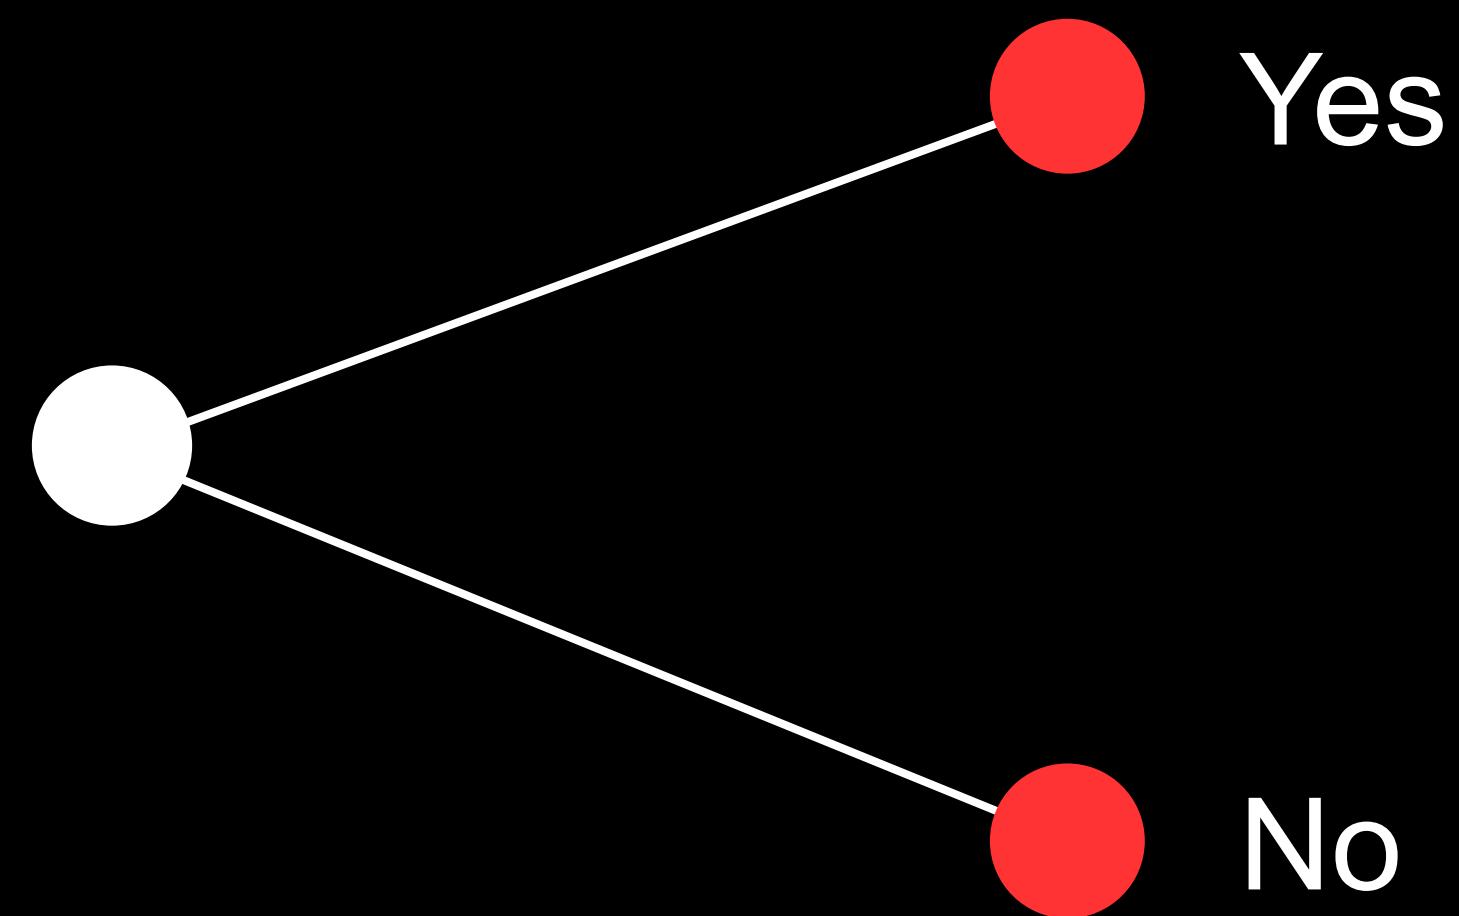


Integration in production



Numerical features

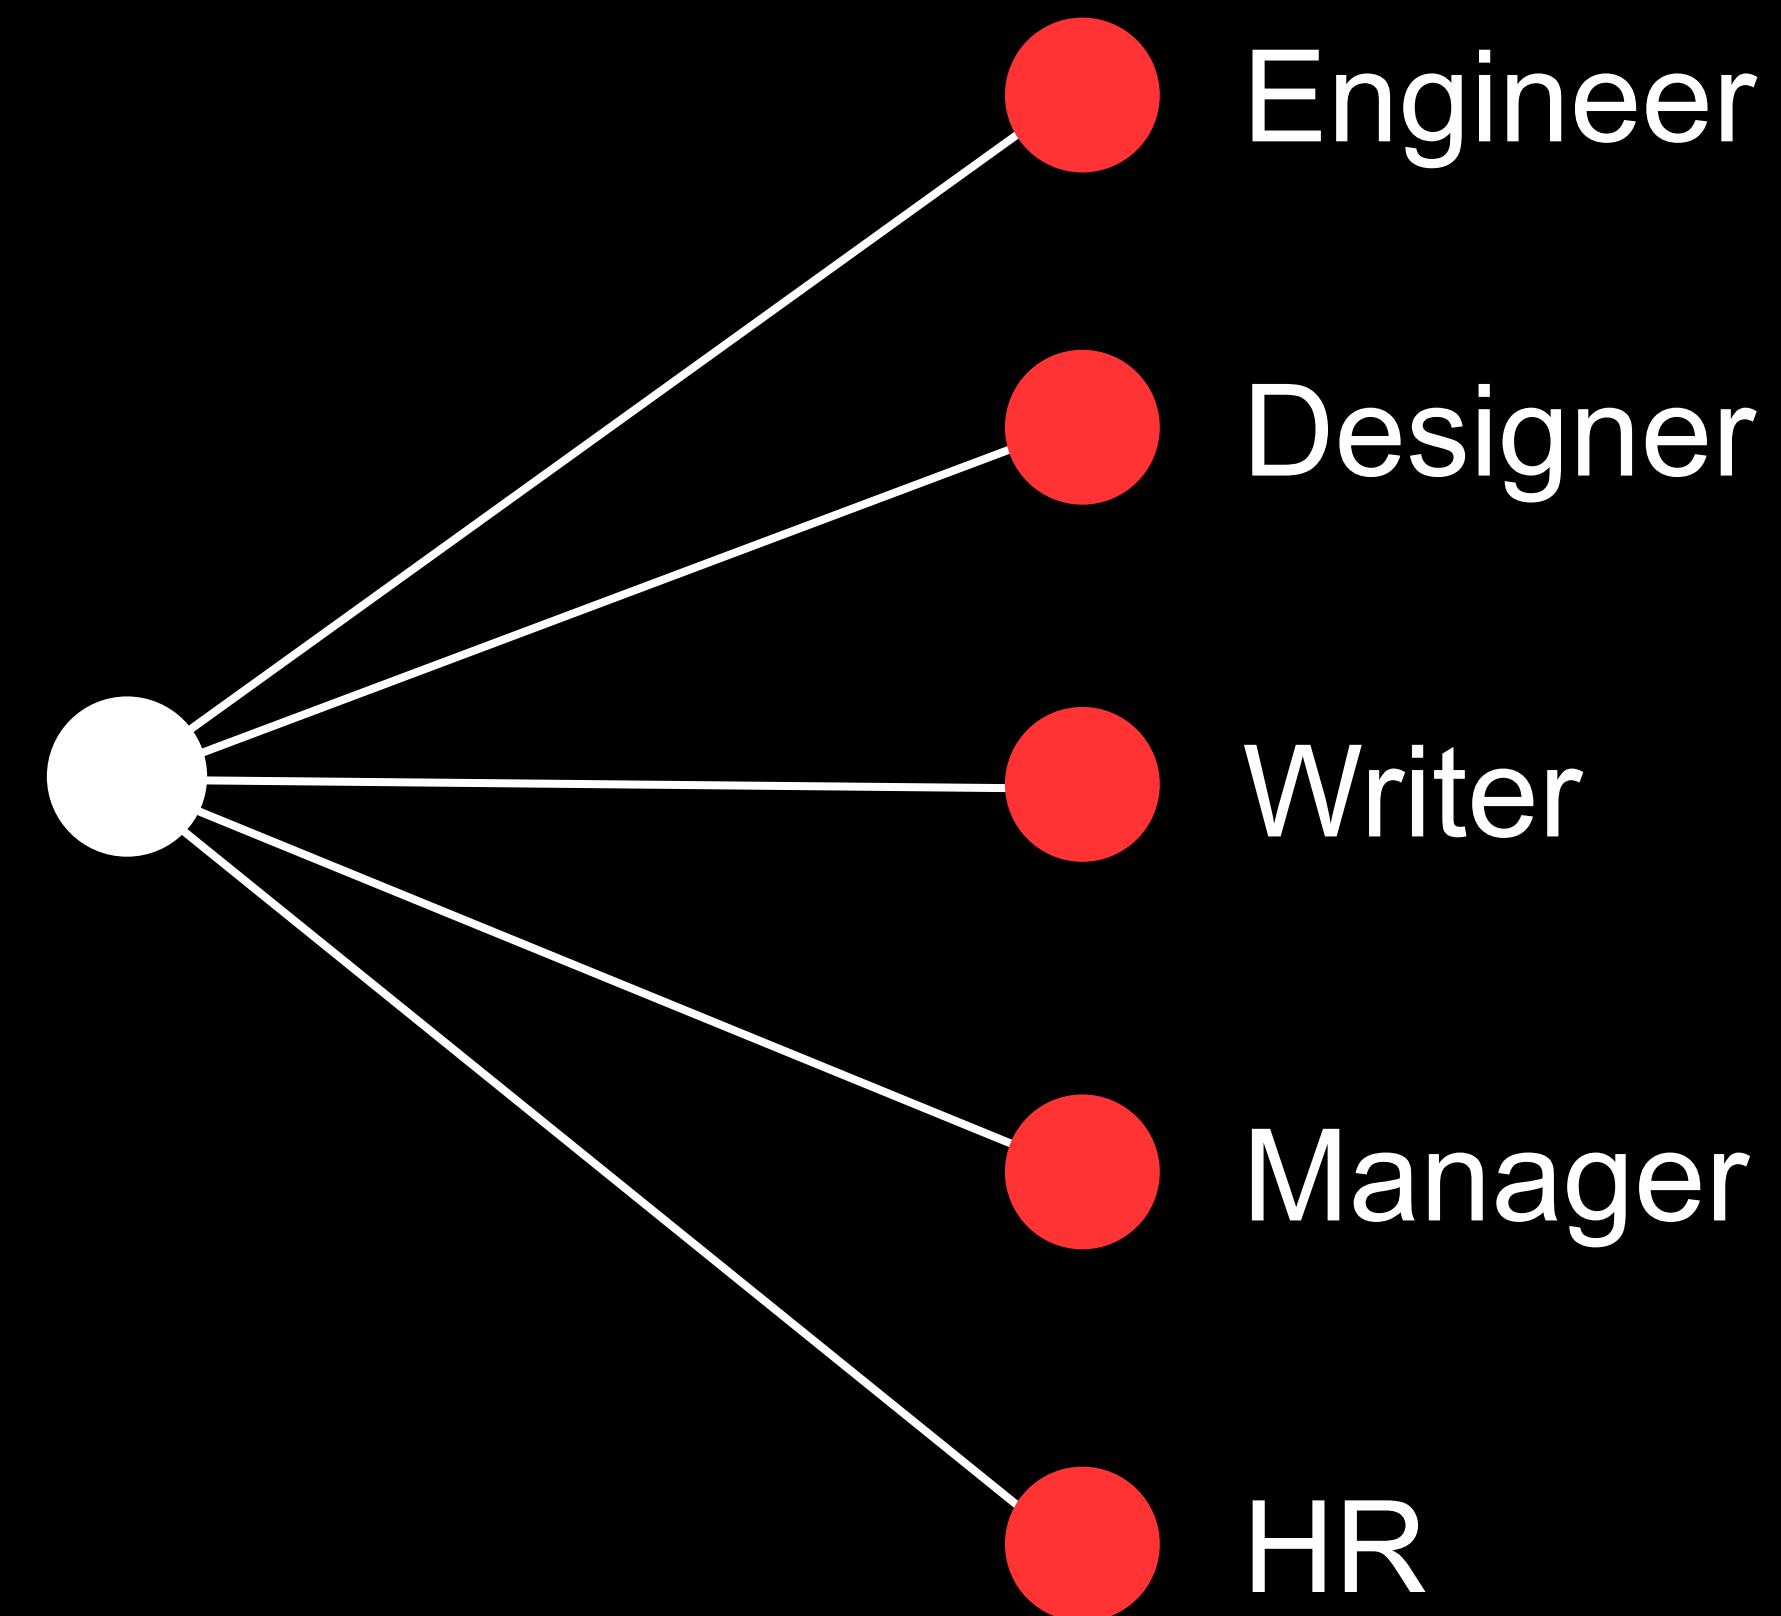
Salary > 30 000



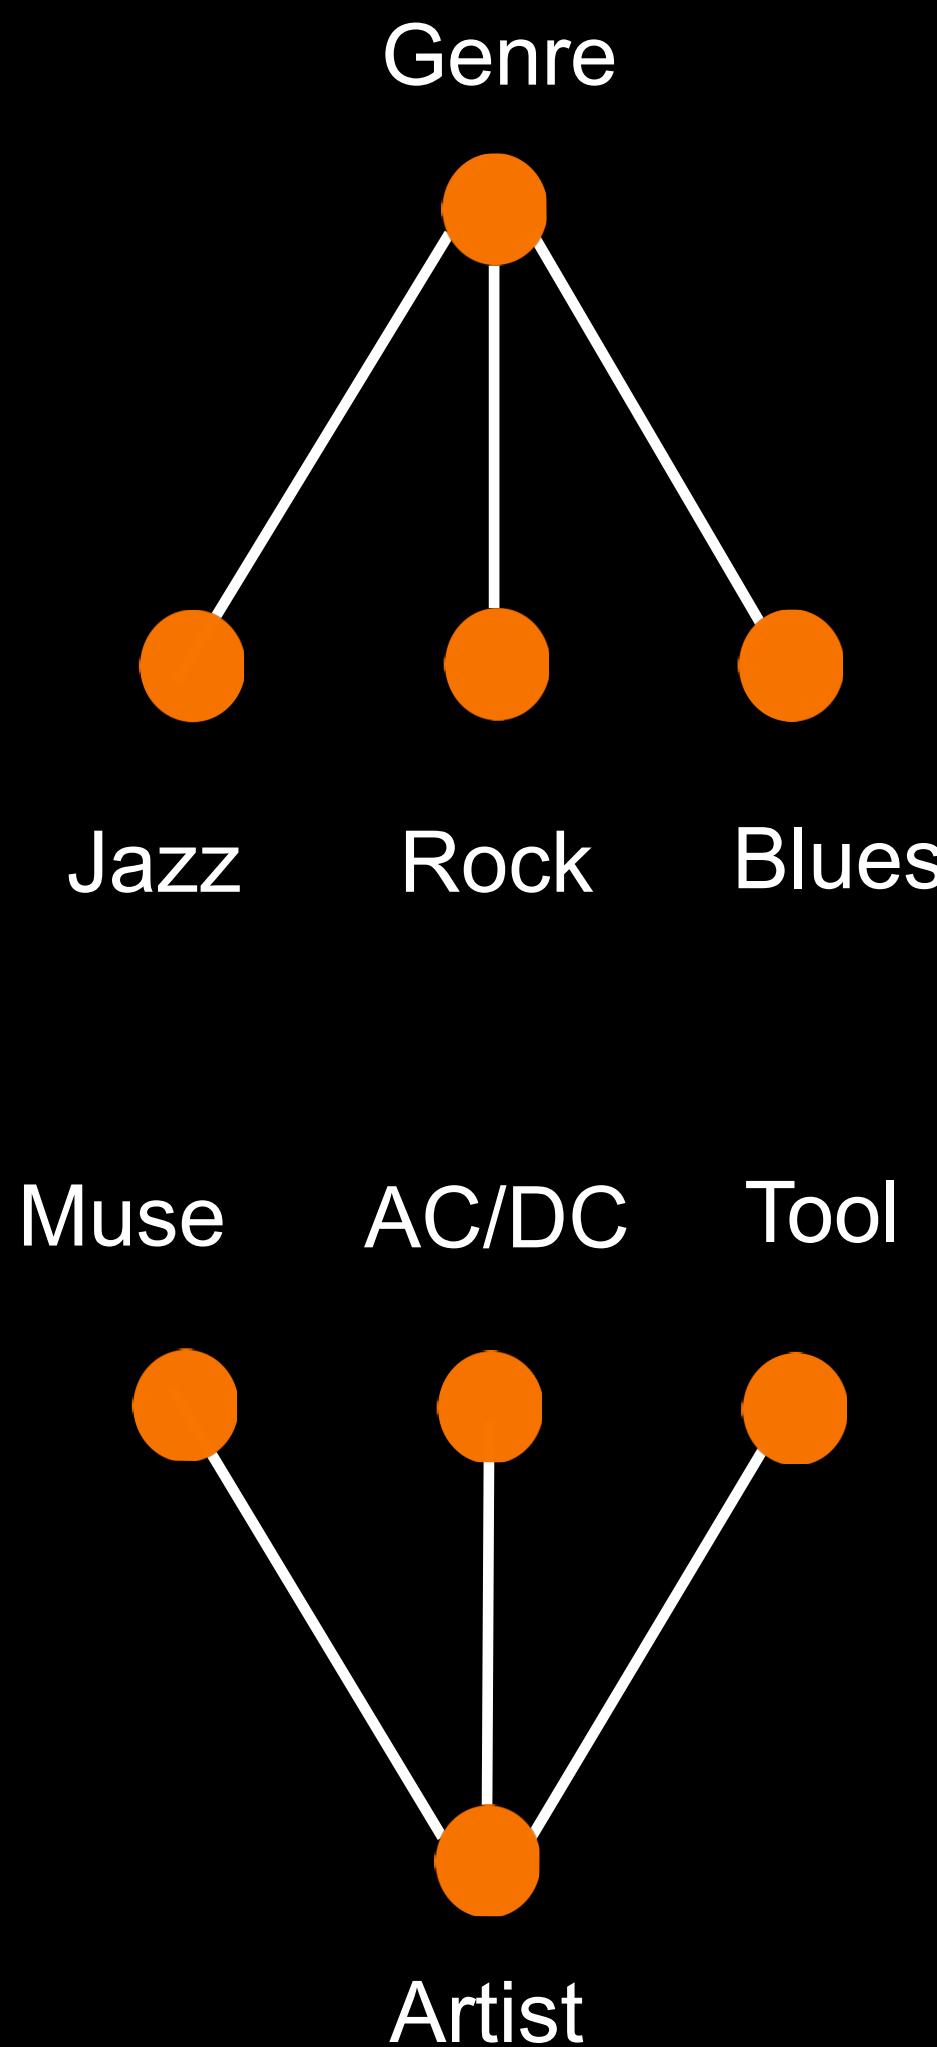
Categorical features

Categorical data

Occupation



Categorical features handling



One-hot encoding

Statistics based on category

Category-based

Label based:
calculated “online”

Greedy search for
combinations

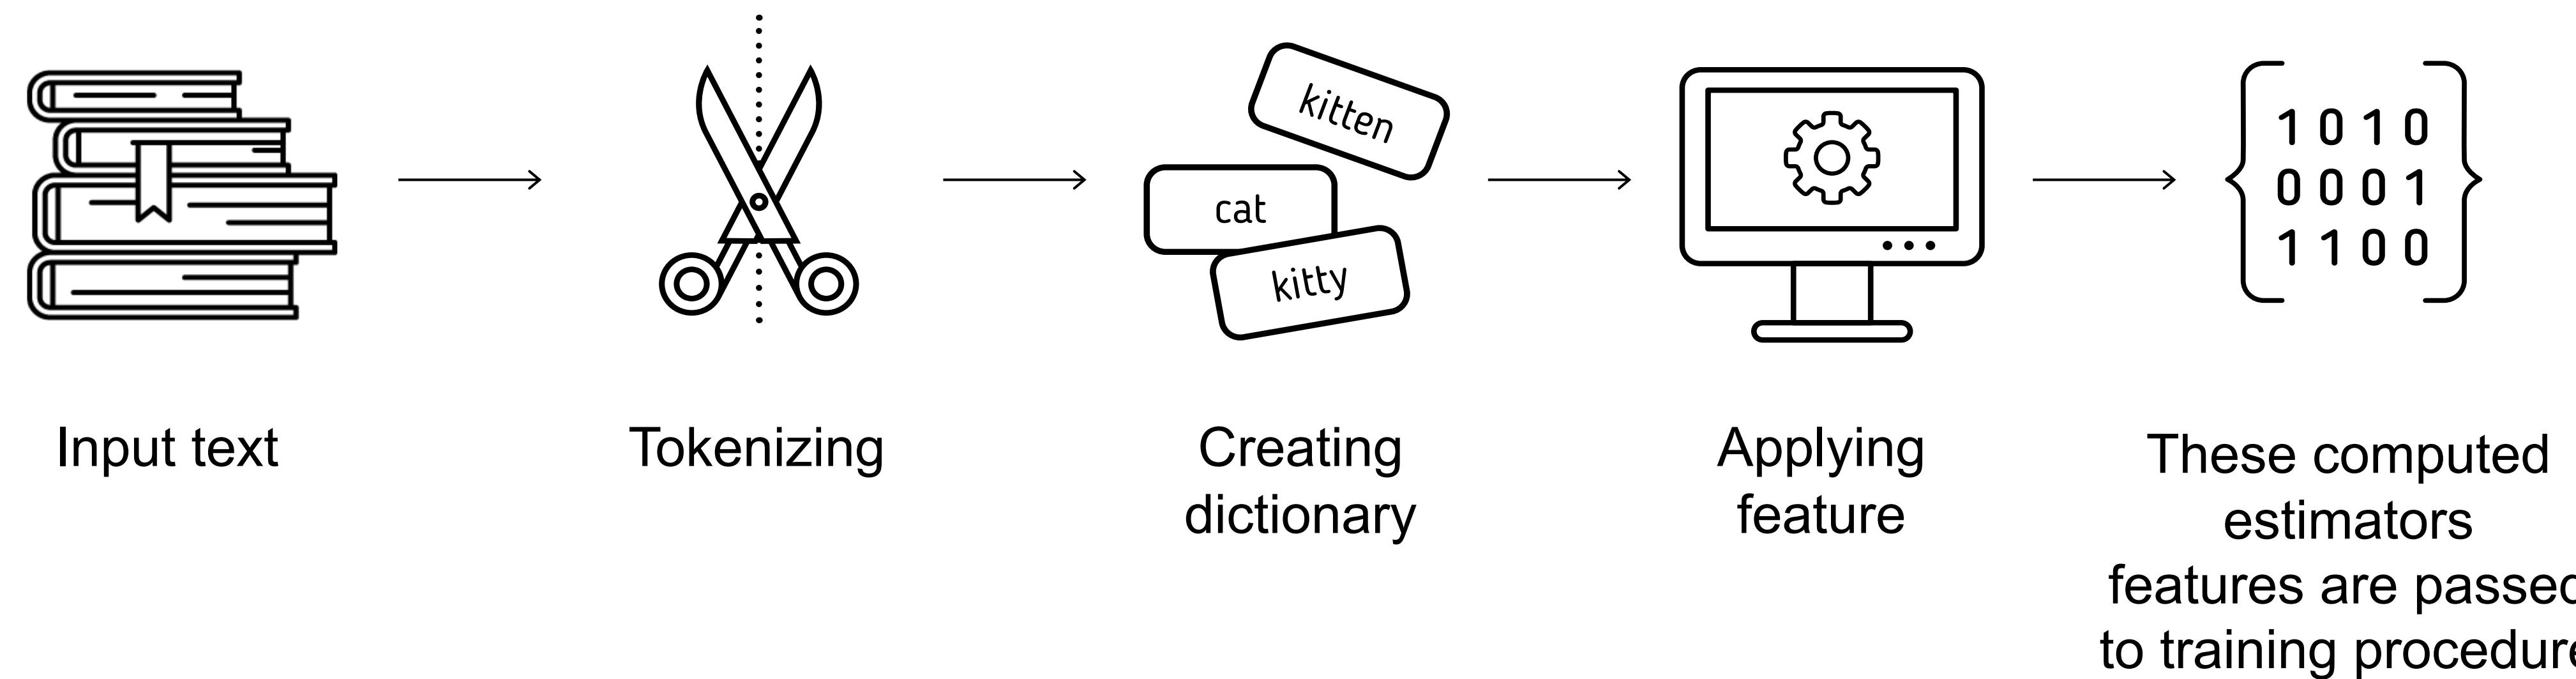
Online features for categories

For every sample “online” feature is calculated using objects with the same category before this one

i	SDE		1
	SDE		1
	SDE		0
	PR		
	SDE		1
	PR		

$$i \longrightarrow \frac{1 + 1 + 0}{3}$$

Text features in CatBoost



Text features

at first i was afraid i
was petrified kept
thinkin i could never
live without you by
my side

[0, 1, 2, 3, 4, 2,
3, 5, 6, 7, 2, 8,
9, 10, 11, 12,
13, 14, 15]



BM25

Naïve Bayes

Bag-of-Words

...

Preprocessing stage

- › Split into words
- › Process numbers and punctuation
- › Build letter and word dictionaries
- › Build ngram dictionaries or BPE



Bag of Words

- › Default: word unigrams + word bigrams
- › For a given dictionary (set of tokens)– is this token present in text?

Dictionary 1:
yesterday, petrified

Dictionary 2:
at+first, i+has,
could+never, ...

at first i was afraid i
was petrified kept
thinkin i could never
live without you by
my side

Feature values:
0,1,1,0,1

Naïve Bayes

- › For every class a new feature $P(\text{Class}|\text{Text})$
- › $P(\text{Class}|\text{Text})$ is replaced with $P(\text{Class}) * \prod_i P(\text{word}_i | \text{Class})$
- › **Most importantly:** calculate it “online”

BM-25 for MultiClass

- › A new numerical feature for every class of multi-classification
- › TF – frequency of a word in text
- › IDF – inverted frequency of a word in a “document”, where “document” is a **concatenation of all texts in this class**
- › **Most importantly**: calculate it “online”

Text processing in CatBoost

- `catboost.Tokenizer`
- `catboost.Dictionary`

Advantages:

- Fast
- Customizable
- Production-ready
- Can be used with other libraries, including Neural Networks

Text features examples

› Rotten Tomatoes: movie review

Numerical

- runtime
- box_office – amount of money raised by ticket sales

Categorical

- critic - name of reviewer
- publisher - journal where the review was published

Text

- review - review of a movie, that was written by a critic
- genres - list of genres that are suitable for this film

review

One very long, dark ride.

genres

Action and Adventure | Art House
and International | Drama |
Mystery and Suspense

Profit from text features

Accuracy on Rotten Tomatoes

Numerical + Categorical
0.4592

+ BOW
0.4616

+ Online Text Features
0.4714

New CatBoost features

New features with effect on quality

- › Automatic learning rate selection for most loss functions
- › Different tree growing strategies both on CPU and GPU
- › Separate quantization for “golden” features
- › Improved MVS sampling
- › New ranking objective StochasticRank
- › Implemented Stochastic Gradient Langevin Boosting
- › Exact calculation for leaf values in some modes

Parameter tuning out of the box

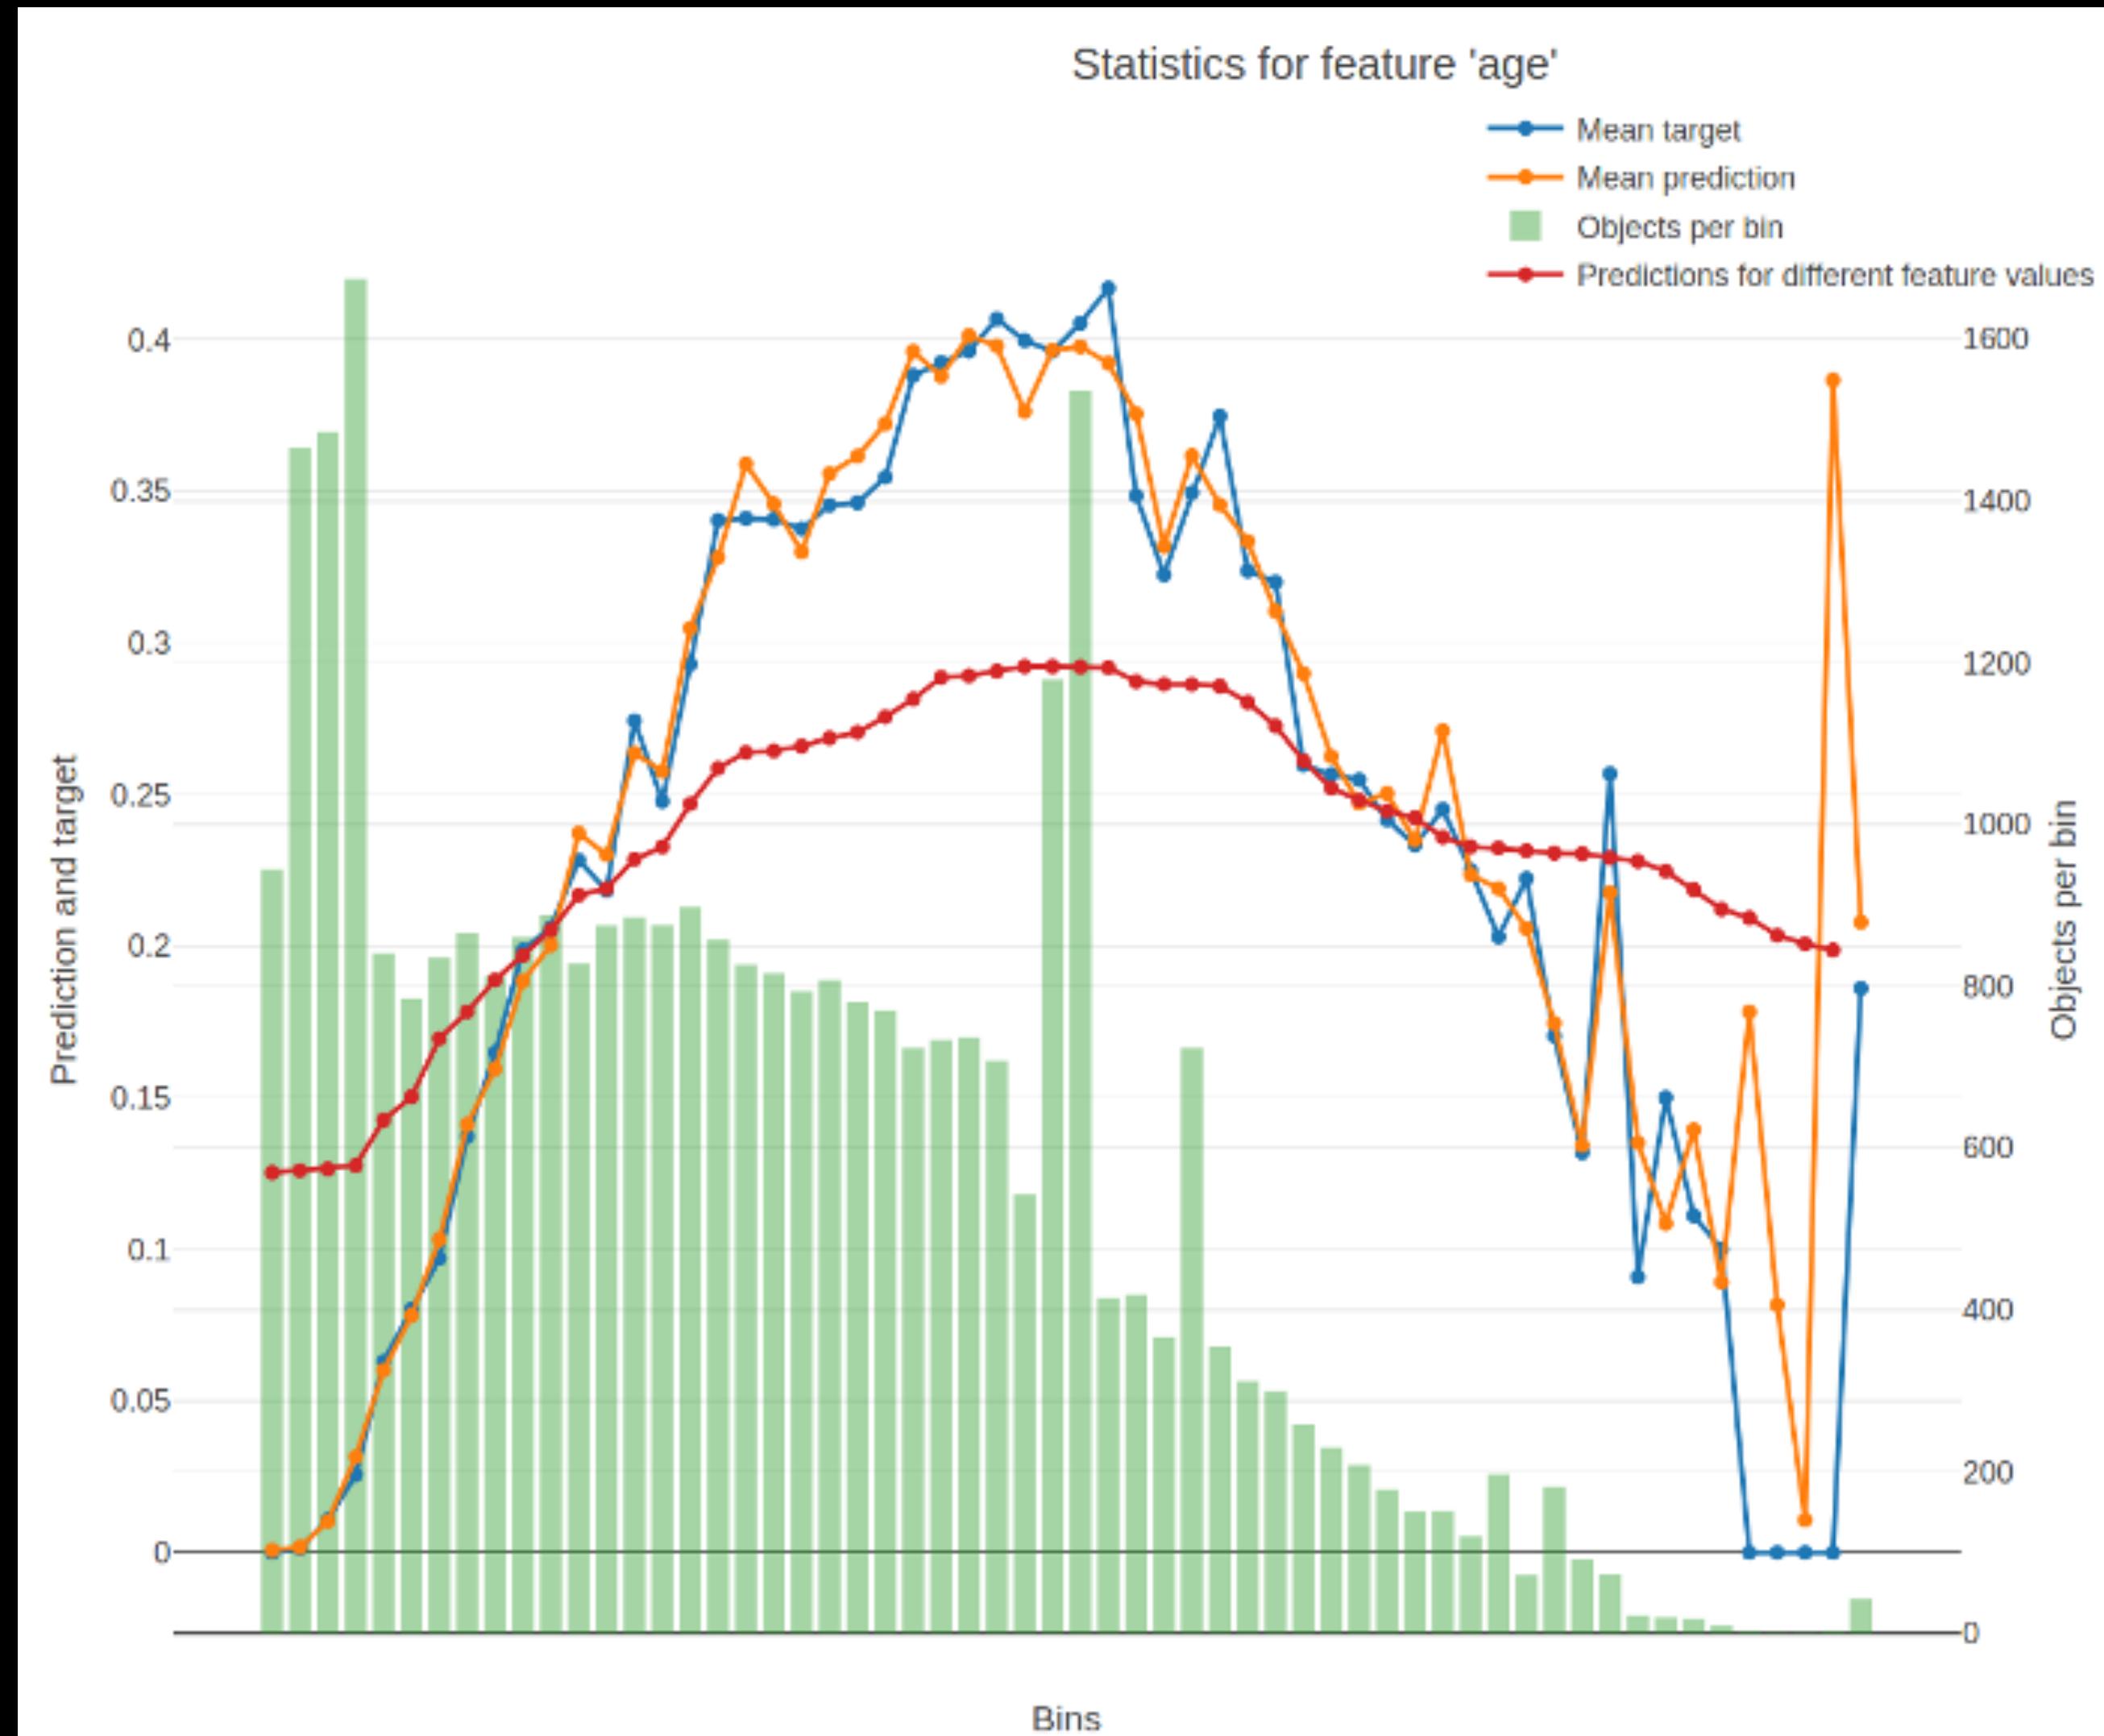
```
In [11]: grid = {  
    'learning_rate': [0.03, 0.1],  
    'depth':[3, 4, 6],  
    'l2_leaf_reg': [1, 3, 5]  
}  
grid_search_results = titanic_model.grid_search(grid, train_pool, shuffle=False, verbose=3, plot=True)
```



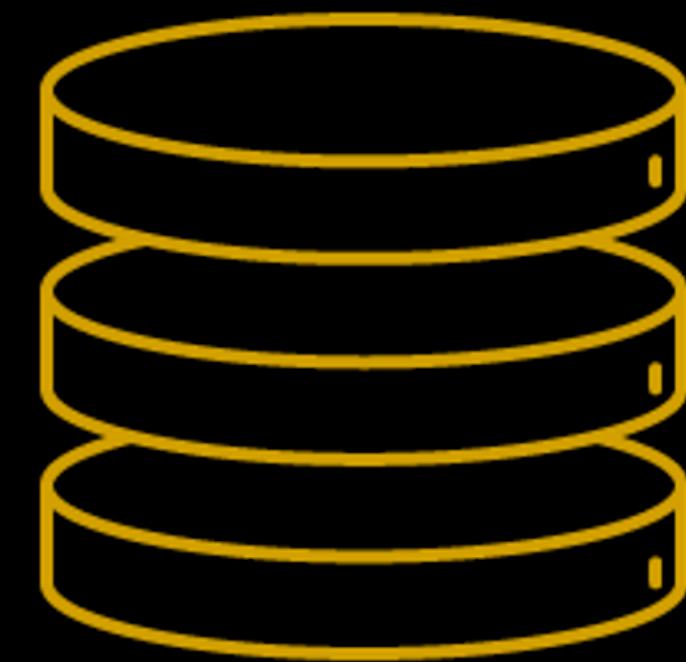
For more information about parameter tuning, visit:
<https://catboost.ai/docs/concepts/parameter-tuning.html>

New Model Analysis tools

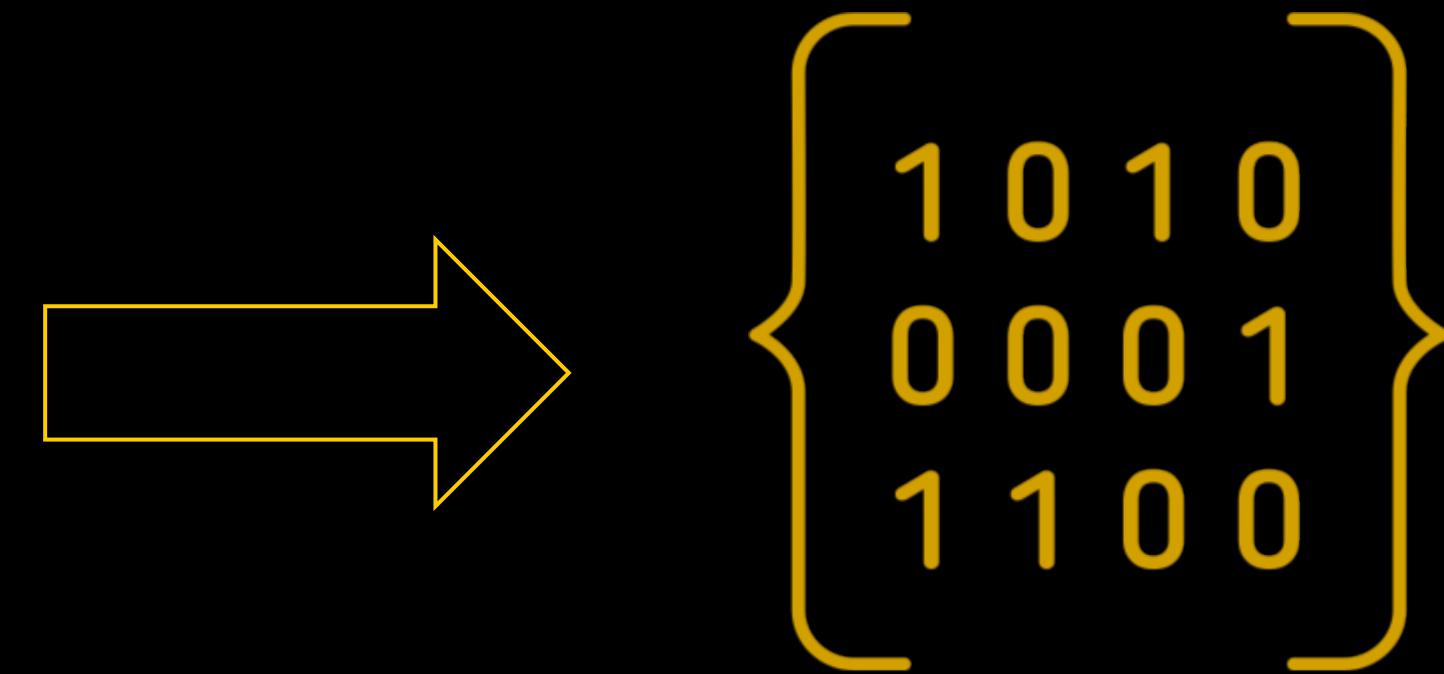
- › Per feature model analysis charts
- › New types of feature importance
- › Tree visualization
- › Ranking analysis



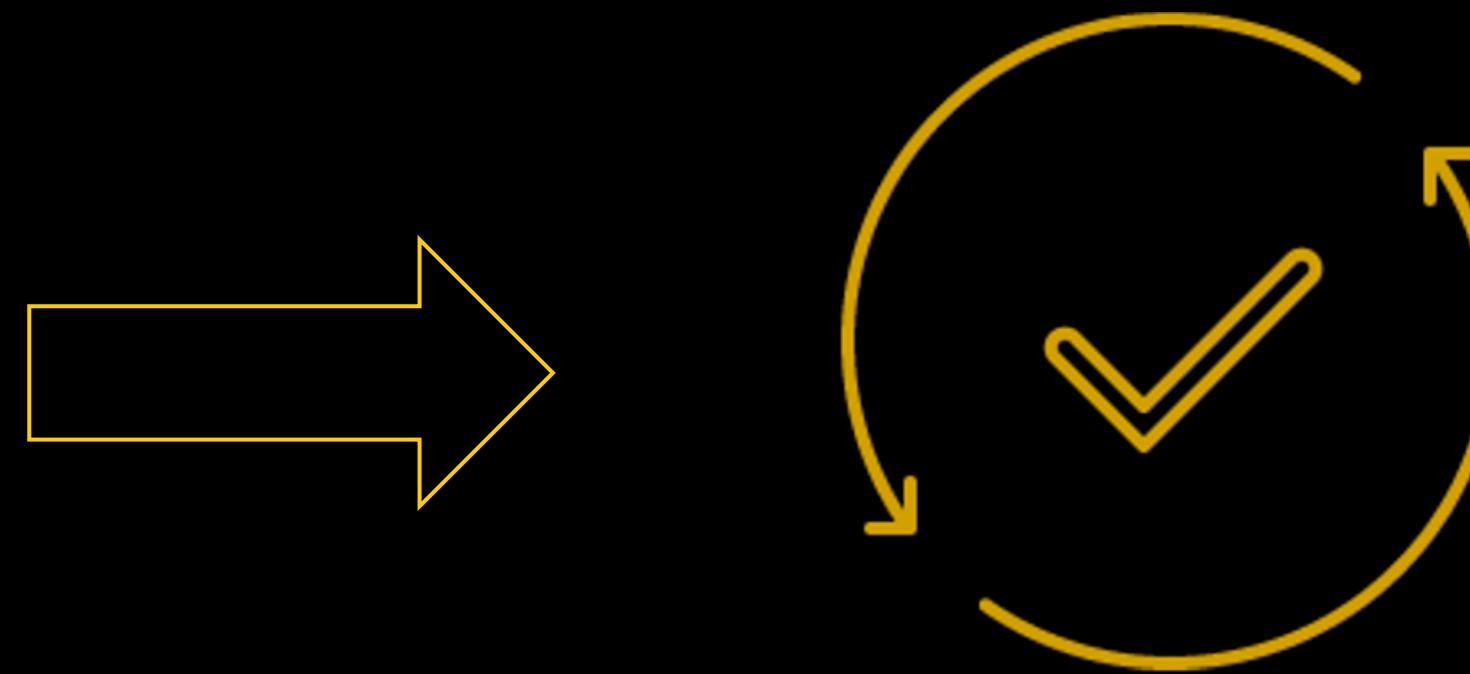
Training on Large Data



Load raw
1TB
(tsv data)



Quantize
~160GB
quantized
representation



Train
~160GB
in RAM or even less on GPU

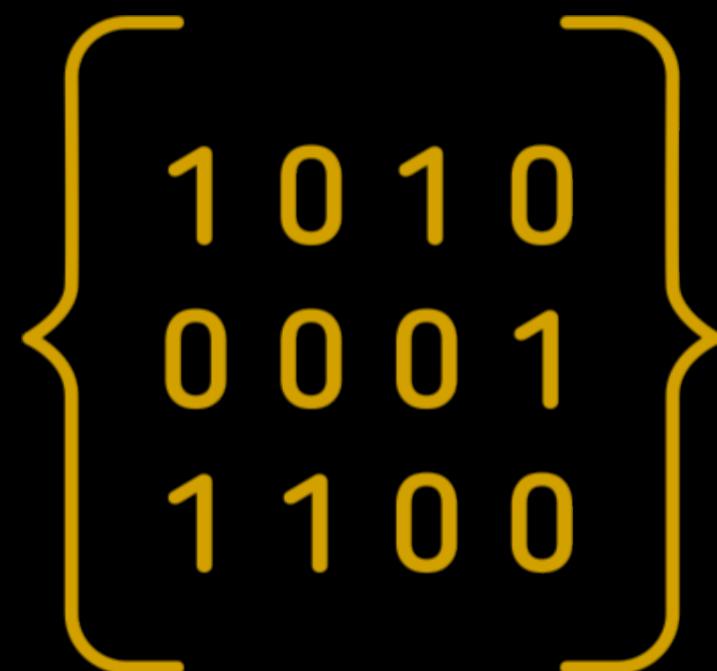
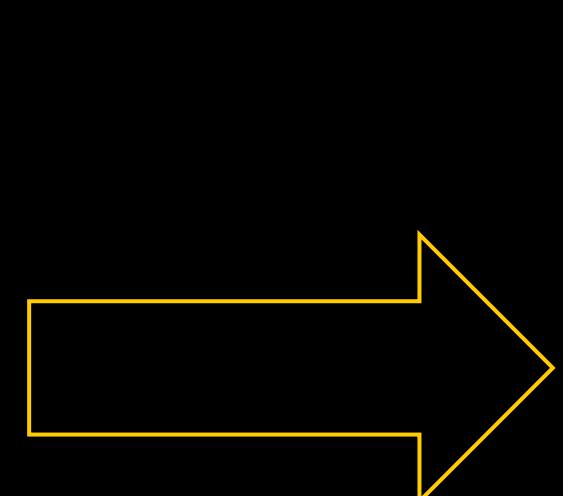
Quantization before training

We've introduced special method `catboost.utils.quantize`
This method quantizes pool in chunks to minimize RAM usage



Load raw

1TB
(tsv data)



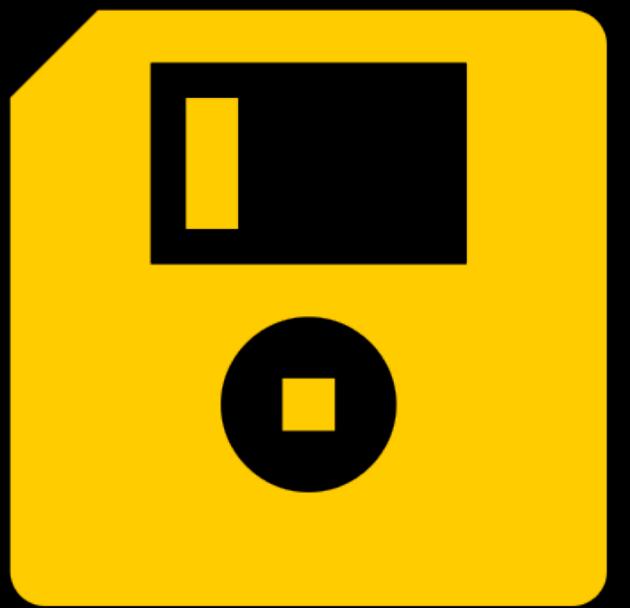
Quantize



Save quantized

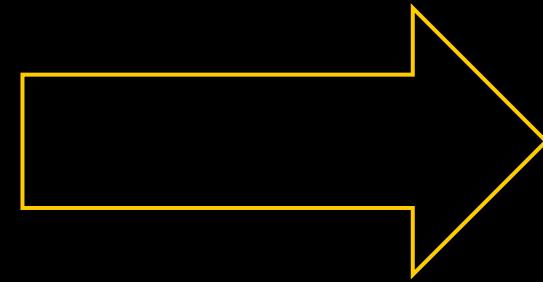
~160GB
quantized
representation

Quantization before training



Load quantized

~160GB
quantized
representation



Train

~160GB
in RAM or even less on GPU

Pool quantization code snippets

```
from catboost import Pool, CatBoostClassifier

# save quantized dataset
train_dataset = Pool(train_data, train_labels)
train_dataset.quantize() # set quantization params here
train_dataset.save('quantized_train.bin')

# we can pass quantized pool path directly to fit
model = CatBoostClassifier()
model.fit('quantized://quantized_train.bin')

# load Pool object from quantized file and use multiple times
quantized_pool = Pool('quantized://quantized_train.bin')
for i in range(N):
    model = CatBoostClassifier()
    model.fit(quantized_pool)
    ...
    ...
```

Pool quantization code snippets

```
from catboost.utils import quantize

pool = quantize(
    data='really_big_train.csv',
    delimiter=',',
    column_description='train.cd',
)
pool.save('big_quantized_train.bin')
```

Speedups

- Huge speedups of preprocessing
- Sparse data support
- Up to 20x speedups of different modes
- Huge speedups for small datasets

Questions?

Kirillov Stanislav

Head of CatBoost team

- › catboost.ai
- › github.com/catboost
- › twitter.com/CatBoostML
- › t.me/catboost_en, t.me/catboost_ru
- › ods.ai => [slack](#) => `tool_catboost` channel